

Improve your workflow for reproducible science

Mine Çetinkaya-Rundel
Duke University + RStudio

 bit.ly/improve-repro-workflow

@minebocek 
mine-cetinkaya-rundel 
cetinkaya.mine@gmail.com 



The results in Table 1
don't seem to correspond to
those in Figure 2!

61

94

45

20

12

44

3

4



```
# set.seed  
set.seed(20190314)
```

```
# generate 8 random numbers between 0 and 99  
runif(8, 0, 99) %>% round()
```

more than

percent

have tried and **failed** to reproduce
another scientist's experiments

more than



percent

have tried and **failed** to reproduce
their **own** experiments

Google Scholar yields



results containing the term **reproducibility crisis**
just in **2021**



earliest reference **reproducibility research***

Jon Claerbout and Martin Karrenbach. "Electronic documents give reproducible research a new meaning."
SEG Technical Program Expanded Abstracts 1992. Society of Exploration Geophysicists, 1992. 601-604.

that I could find...

SUMMARY

A revolution in education and technology transfer follows from the marriage of word processing and software command scripts. In this marriage an author attaches to every figure caption a pushbutton or a name tag usable to recalculate the figure from all its data, parameters, and programs. This provides a concrete definition of reproducibility in computationally oriented research. Experience at the Stanford Exploration Project shows that preparing such electronic documents is little effort beyond our customary report writing; mainly, we need to file everything in a systematic way.

In 1990 we began experimenting with electronic documents that merge our scientific software with our word-processing software. A year later we manufactured a CD-ROM containing a new textbook, Joe Dellinger's doctoral dissertation, and two progress reports of the Stanford Exploration Project. We distributed these CD-ROMs¹ to sponsors and many friends at the 1991 SEG meeting.

In 1990, we set this sequence of goals:

- Learn how to merge a publication with its underlying computational analysis.
- Teach researchers how to prepare a document in a form where they themselves can reproduce their own research results a year or more later by "pressing a single button".
- Learn how to leave finished work in a condition where coworkers can reproduce the calculation including the final illustration by pressing a button in its caption.
- Prepare a complete copy of our local software environment so that graduating students can take their work away with them to other sites, press a button, and reproduce their Stanford work.
- Merge electronic documents written by multiple authors (SEP reports).
- Export electronic documents to numerous other sites (sponsors) so they can readily reproduce a substantial portion of our Stanford research.

We met all these goals and set new ones:

- produce all new documents in this form, including lab reports in formal classes and "lab notebooks" of research progress.

¹SEP-CD-I is available from Stanford University Press, \$15 plus shipping, tel 415-723-1593

- make incremental improvements in electronic-document software
- seek partners for broadening standards (and making incremental improvements).

Our basic goal is reproducible research. The electronic document is our means to this end. In principle, reproducibility in research can be achieved without electronic documents and that is how we started. Our first nonelectronic reproducible document was a textbook in which the paper document contained the name of a program script in every figure caption. The program scripts were organized by book chapter and section so they could be correlated to an accompanying magnetic tape dump of the file system. The magnetic tape also contained all the necessary data to feed the program script.

Now that we have begun using CD-ROM publication, we can go much further. Every figure caption contains a pushbutton that jumps to the appropriate science directory (folder) and initiates a figure rebuild command and then displays the figure, possibly as a movie or interactive program. We normally display seismic images of the earth's interior, but to reach wider audiences, Figure 1 shows a satellite weather picture which the pushbutton will animate as seen on commercial television. We include all our plot software as well as freely available software from many sources, including compilers and the \LaTeX word processing system. Naturally we cannot include licensed software, but with the exception of Fortran and C compilers and the UNIX system itself, our publication includes source code for everything needed. The CD-ROM, at 680 megabytes, is so large we have had room for many executable programs on popular brands of workstations. The presence of these executables gives our readers a fast start.

Nearly everyone would rather read a paper book than the bitmapped page images on a screen that you see with an electronic document. But the illustrations in the electronic book are mostly in color, many are movies, and some are interactive. So the electronic book gives the reader a better understanding of the results. We typically use an interactive movie program to compare seismic sections where successive frames include processing with various parameters. The movie medium is much more informative than comparing seismic sections side by side. 3-D volumes are much better exhibited by movies than static paper illustrations. We are delivering a volume of software that is accessed like a book.

In 1990, we set this sequence of goals:

- Learn how to merge a publication with its underlying computational analysis.
- Teach researchers how to prepare a document in a form where they themselves can reproduce their own research results a year or more later by "pressing a single button".
- Learn how to leave finished work in a condition where coworkers can reproduce the calculation including the final illustration by pressing a button in its caption.
- Prepare a complete copy of our local software environment so that graduating students can take their work away with them to other sites, press a button, and reproduce their Stanford work.
- Merge electronic documents written by multiple authors (SEP reports).
- Export electronic documents to numerous other sites (sponsors) so they can readily reproduce a substantial portion of our Stanford research.

SUMMARY

A revolution in education and technology transfer follows from the marriage of word processing and software command scripts. In this marriage an author attaches to every figure caption a pushbutton or a name tag usable to recalculate the figure from all its data, parameters, and programs. This provides a concrete definition of reproducibility in computationally oriented research. Experience at the Stanford Exploration Project shows that preparing such electronic documents is little effort beyond our customary report writing; mainly, we need to file everything in a systematic way.

In 1990 we began experimenting with electronic documents that merge our scientific software with our word-processing software. A year later we manufactured a CD-ROM containing a new textbook, Joe Dellinger's doctoral dissertation, and two progress reports of the Stanford Exploration Project. We distributed these CD-ROMs¹ to sponsors and many friends at the 1991 SEG meeting.

In 1990, we set this sequence of goals:

- Learn how to merge a publication with its underlying computational analysis.
- Teach researchers how to prepare a document in a form where they themselves can reproduce their own research results a year or more later by "pressing a single button".
- Learn how to leave finished work in a condition where coworkers can reproduce the calculation including the final illustration by pressing a button in its caption.
- Prepare a complete copy of our local software environment so that graduating students can take their work away with them to other sites, press a button, and reproduce their Stanford work.
- Merge electronic documents written by multiple authors (SEP reports).
- Export electronic documents to numerous other sites (sponsors) so they can readily reproduce a substantial portion of our Stanford research.

We met all these goals and set new ones:

- produce all new documents in this form, including lab reports in formal classes and "lab notebooks" of research progress.

- make incremental improvements in electronic-document software
- seek partners for broadening standards (and making incremental improvements).

Our basic goal is reproducible research. The electronic document is our means to this end. In principle, reproducibility in research can be achieved without electronic documents and that is how we started. Our first nonelectronic reproducible document was a textbook in which the paper document contained the name of a program script in every figure caption. The program scripts were organized by book chapter and section so they could be correlated to an accompanying magnetic tape dump of the file system. The magnetic tape also contained all the necessary data to feed the program script.

Now that we have begun using CD-ROM publication, we can go much further. Every figure caption contains a pushbutton that jumps to the appropriate science directory (folder) and initiates a figure rebuild command and then displays the figure, possibly as a movie or interactive program. We normally display seismic images of the earth's interior, but to reach wider audiences, Figure 1 shows a satellite weather picture which the pushbutton will animate as seen on commercial television. We include all our plot software as well as freely available software from many sources, including compilers and the L^AT_EX word processing system. Naturally we cannot include licensed software, but with the exception of Fortran and C compilers and the UNIX system itself, our publication includes source code for everything needed. The CD-ROM, at 680 megabytes, is so large we have had room for many executable programs on popular brands of workstations. The presence of these executables gives our readers a fast start.

Nearly everyone would rather read a paper book than the bitmapped page images on a screen that you see with an electronic document. But the illustrations in the electronic book are mostly in color, many are movies, and some are interactive. So the electronic book gives the reader a better understanding of the results. We typically use an interactive movie program to compare seismic sections where successive frames include processing with various parameters. The movie medium is much more informative than comparing seismic sections side by side. 3-D volumes are much better exhibited by movies than static paper illustrations. We are delivering a volume of software that is accessed like a book.

¹SEP-CD-1 is available from Stanford University Press, \$15 plus shipping, tel 415-723-1593

Now that we have begun using CD-ROM publication, we can go much further. Every figure caption contains a pushbutton that jumps to the appropriate science directory (folder) and initiates a figure rebuild command and then displays the figure, possibly as a movie or interactive program. We normally display seismic images of the earth's inter-

Pioneering 'live-code' article allows scientists to play with each other's results

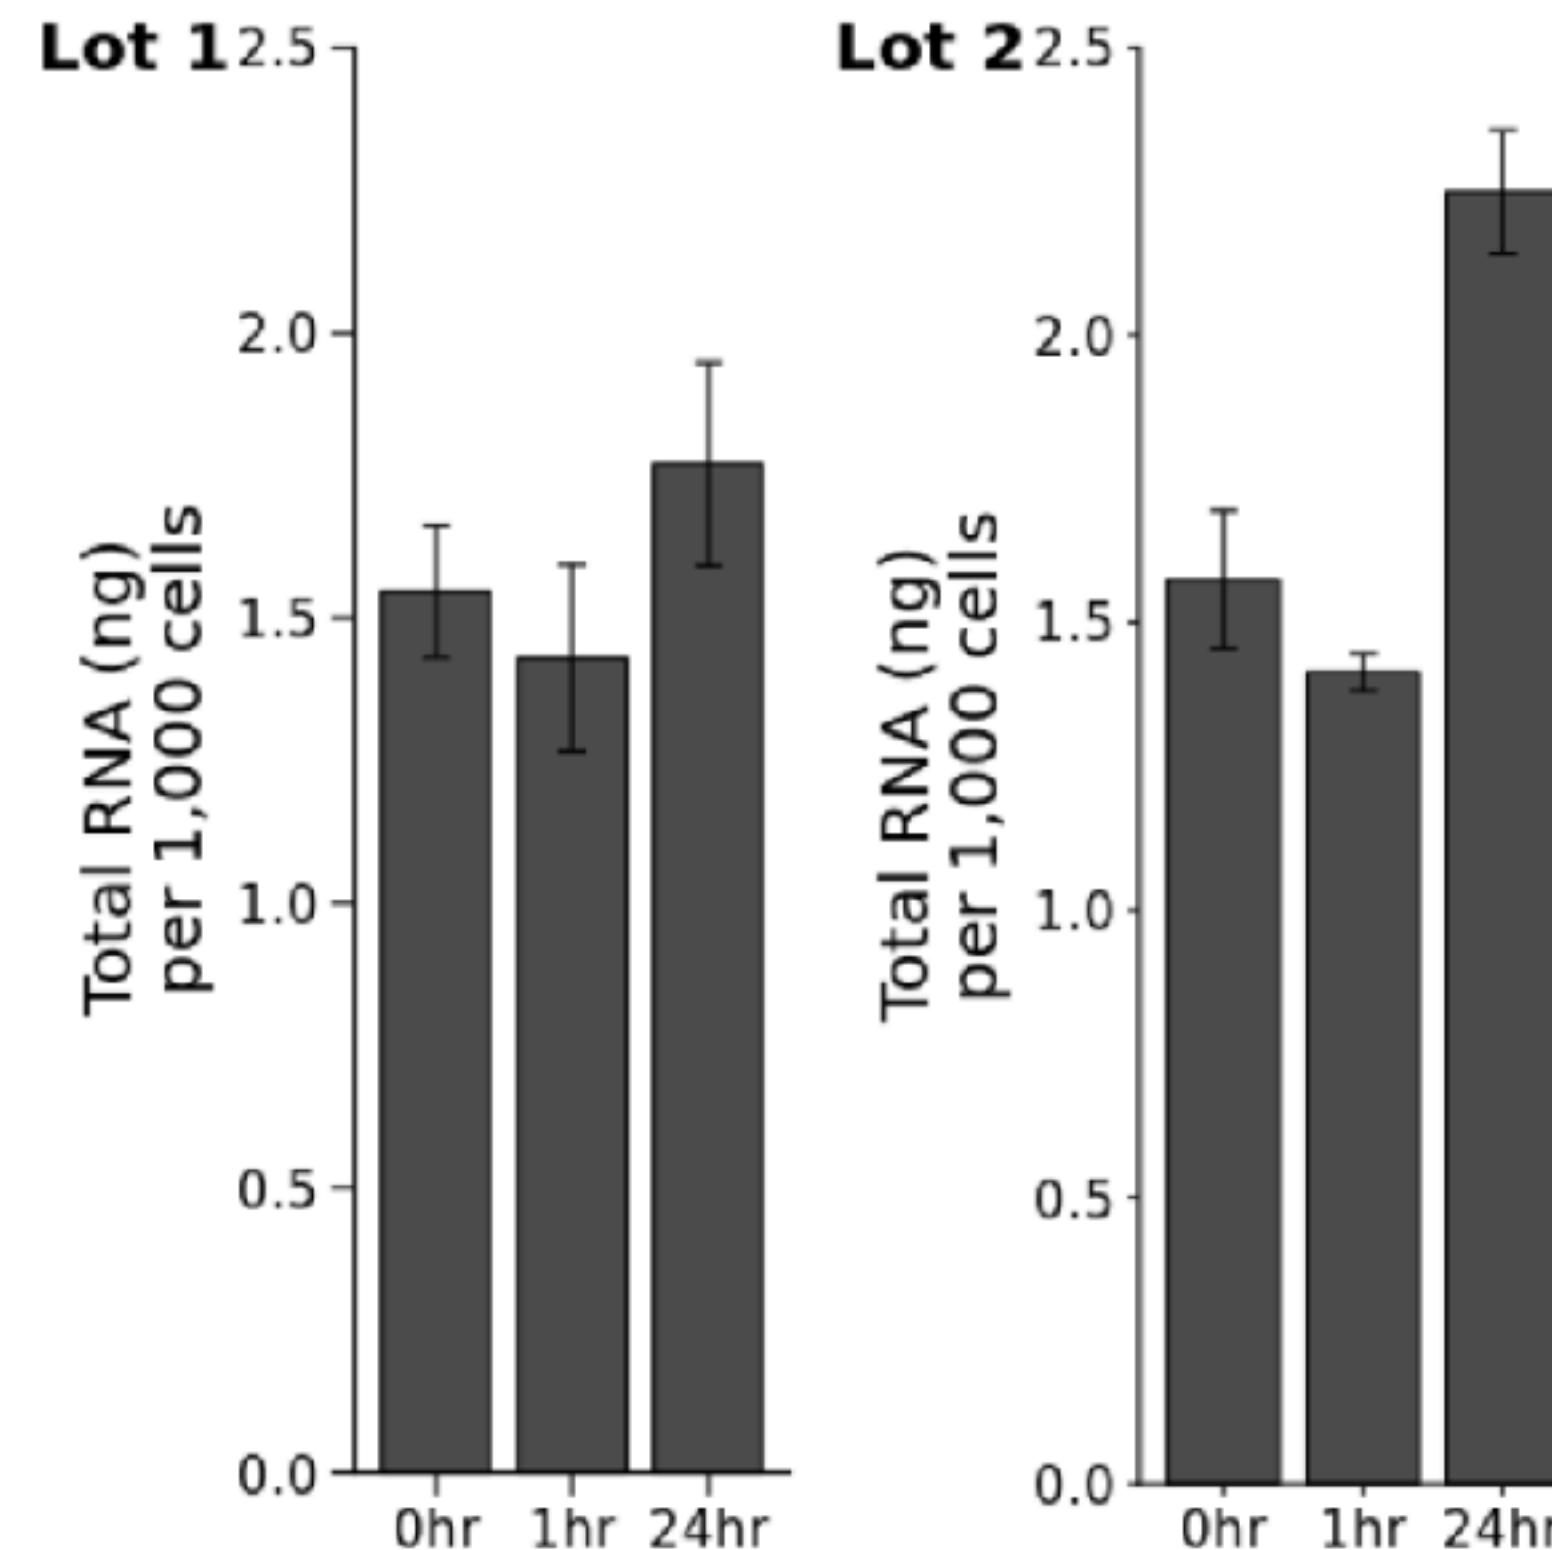
eLife's prototype lets scientists *modify the software underlying figures to validate, build on, or better understand the work.*

This is a [Reproducible document](#). See the [original article or source](#).

NOTE: Below is a reproducible version of Figure 1B. You can inspect the code, make changes and run the code by pressing SHIFT+ENTER. The data used can be [downloaded here](#).

R Script

status: **ready** (run code with ↵)



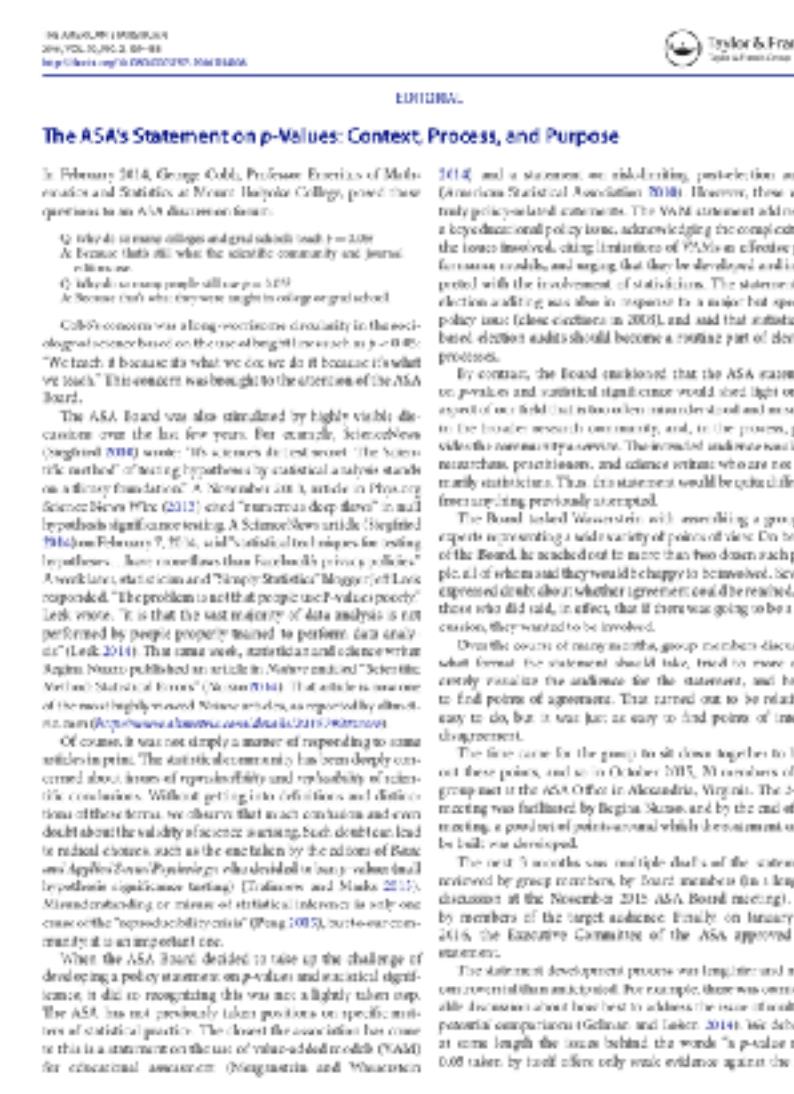
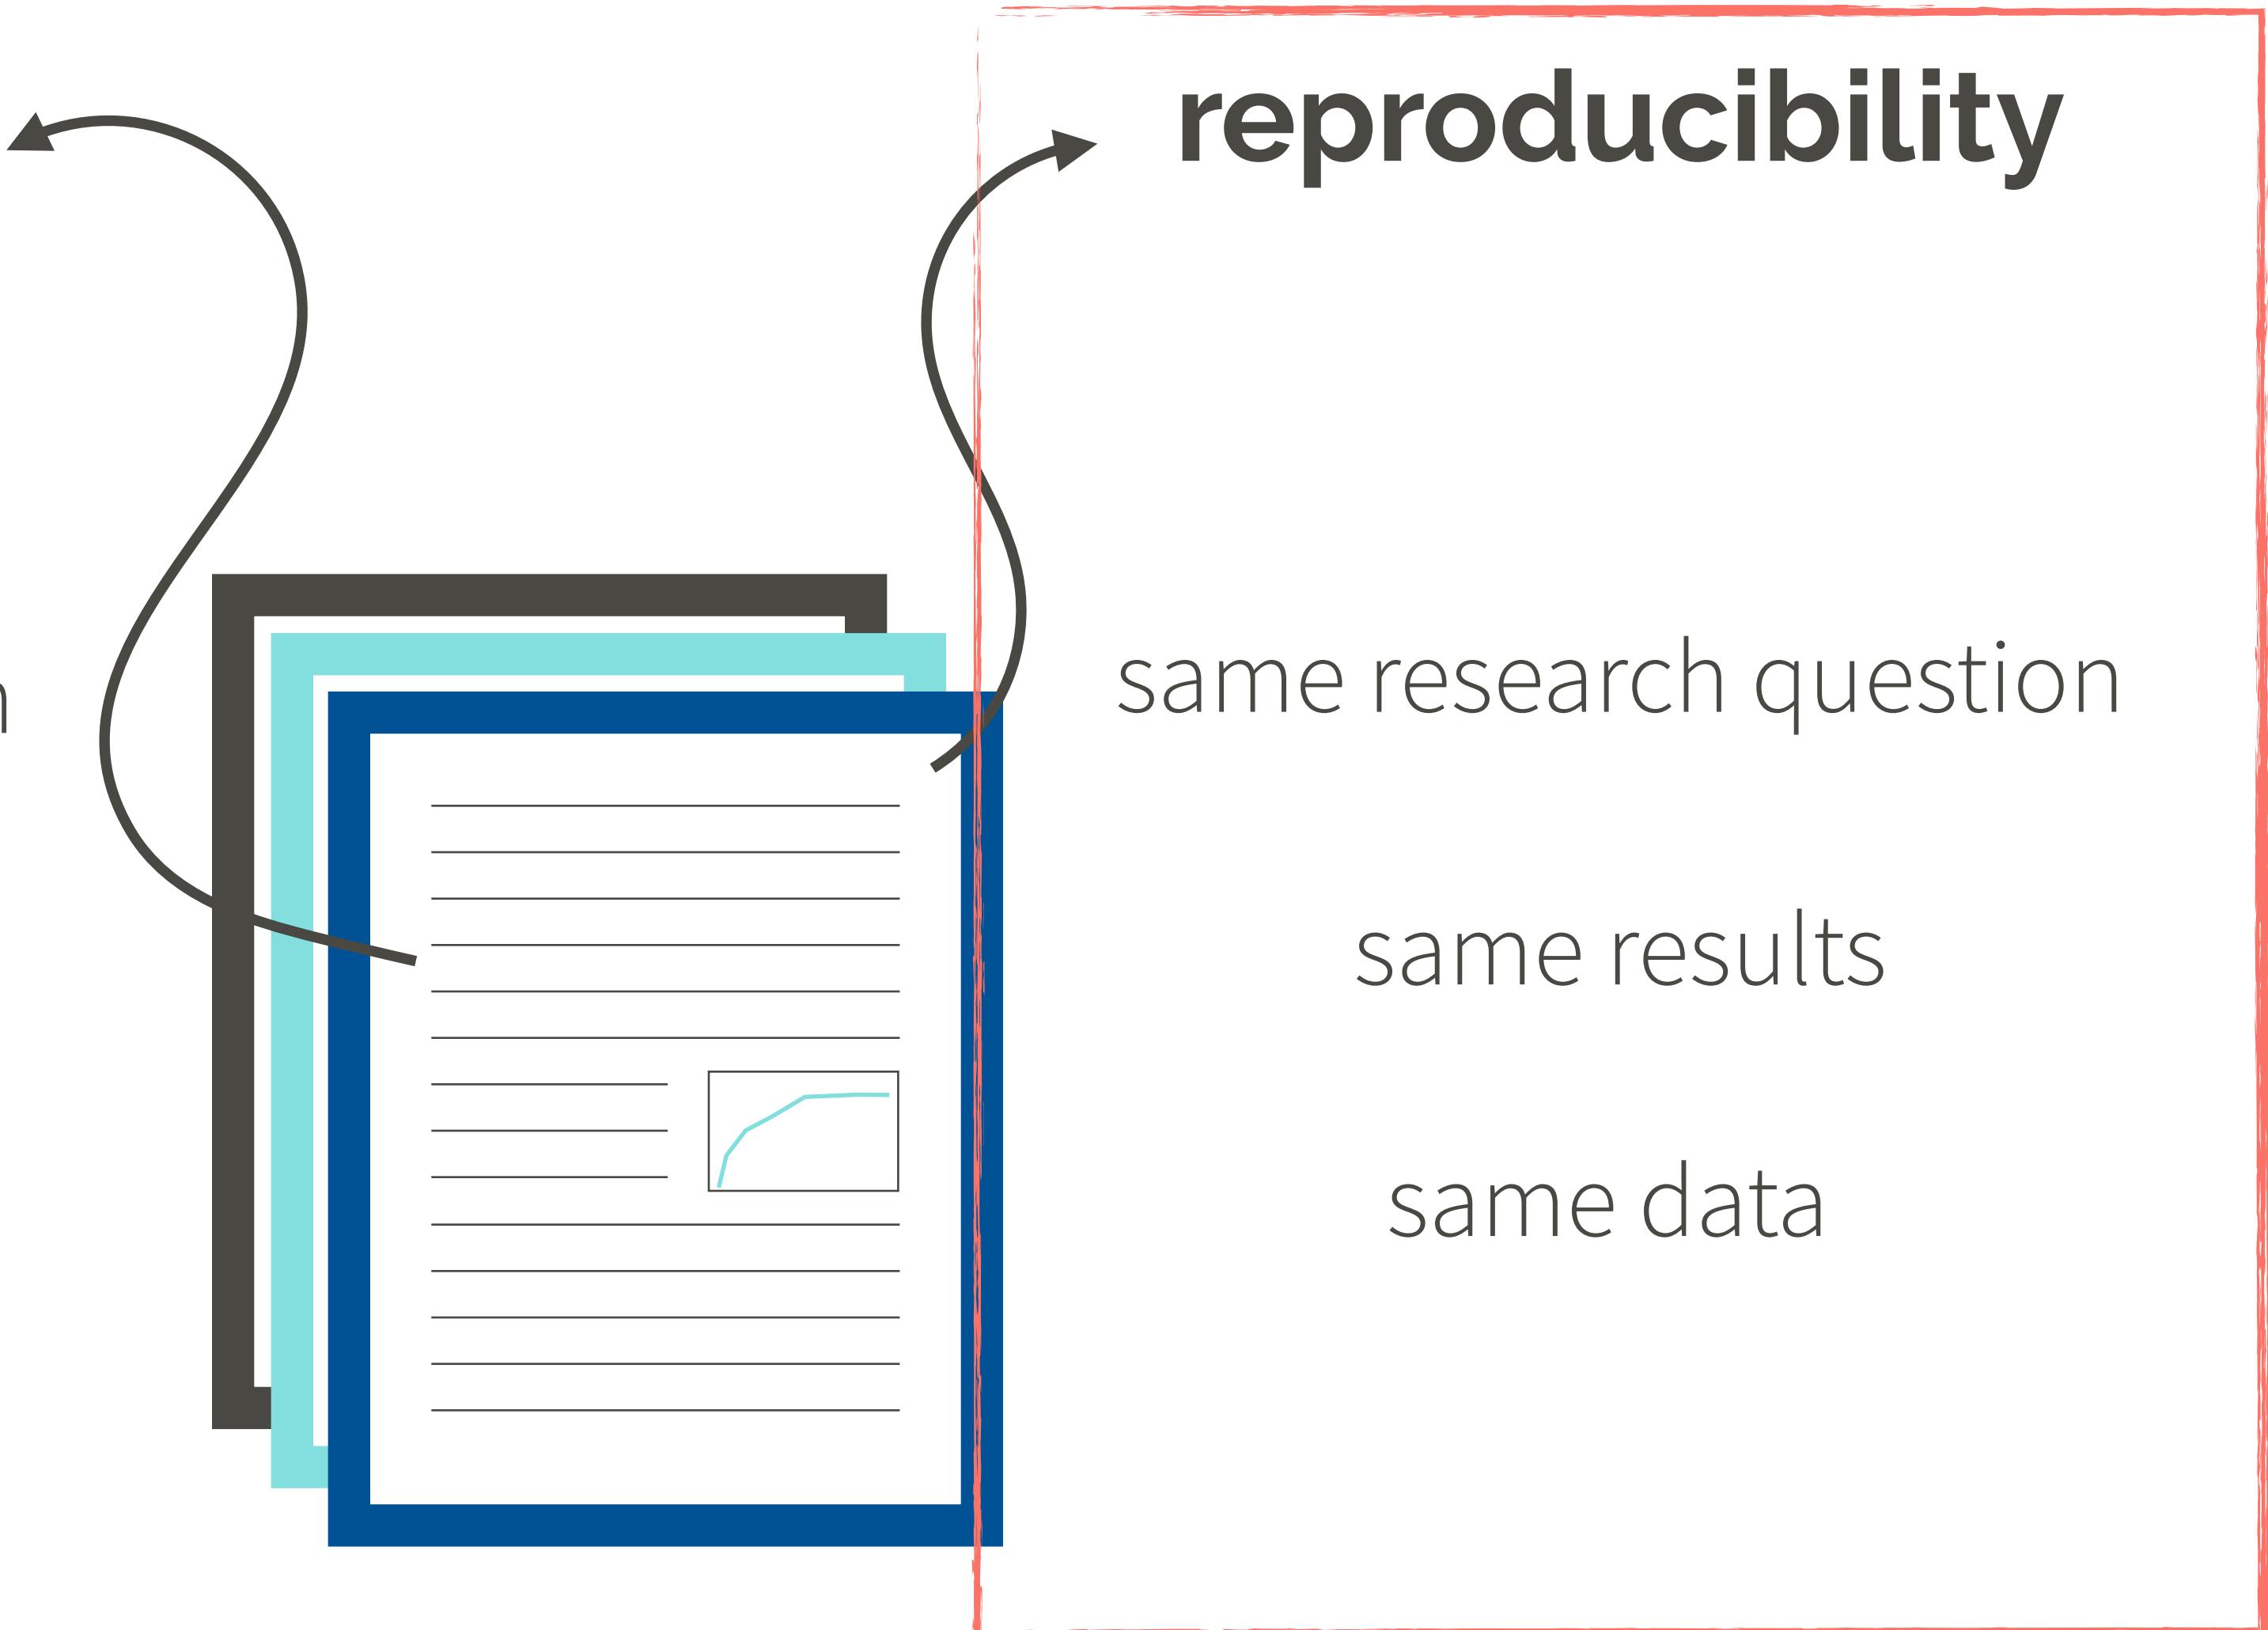
Total RNA levels following c-Myc overexpression



setting the stage

replicability

same research question
same results
new data

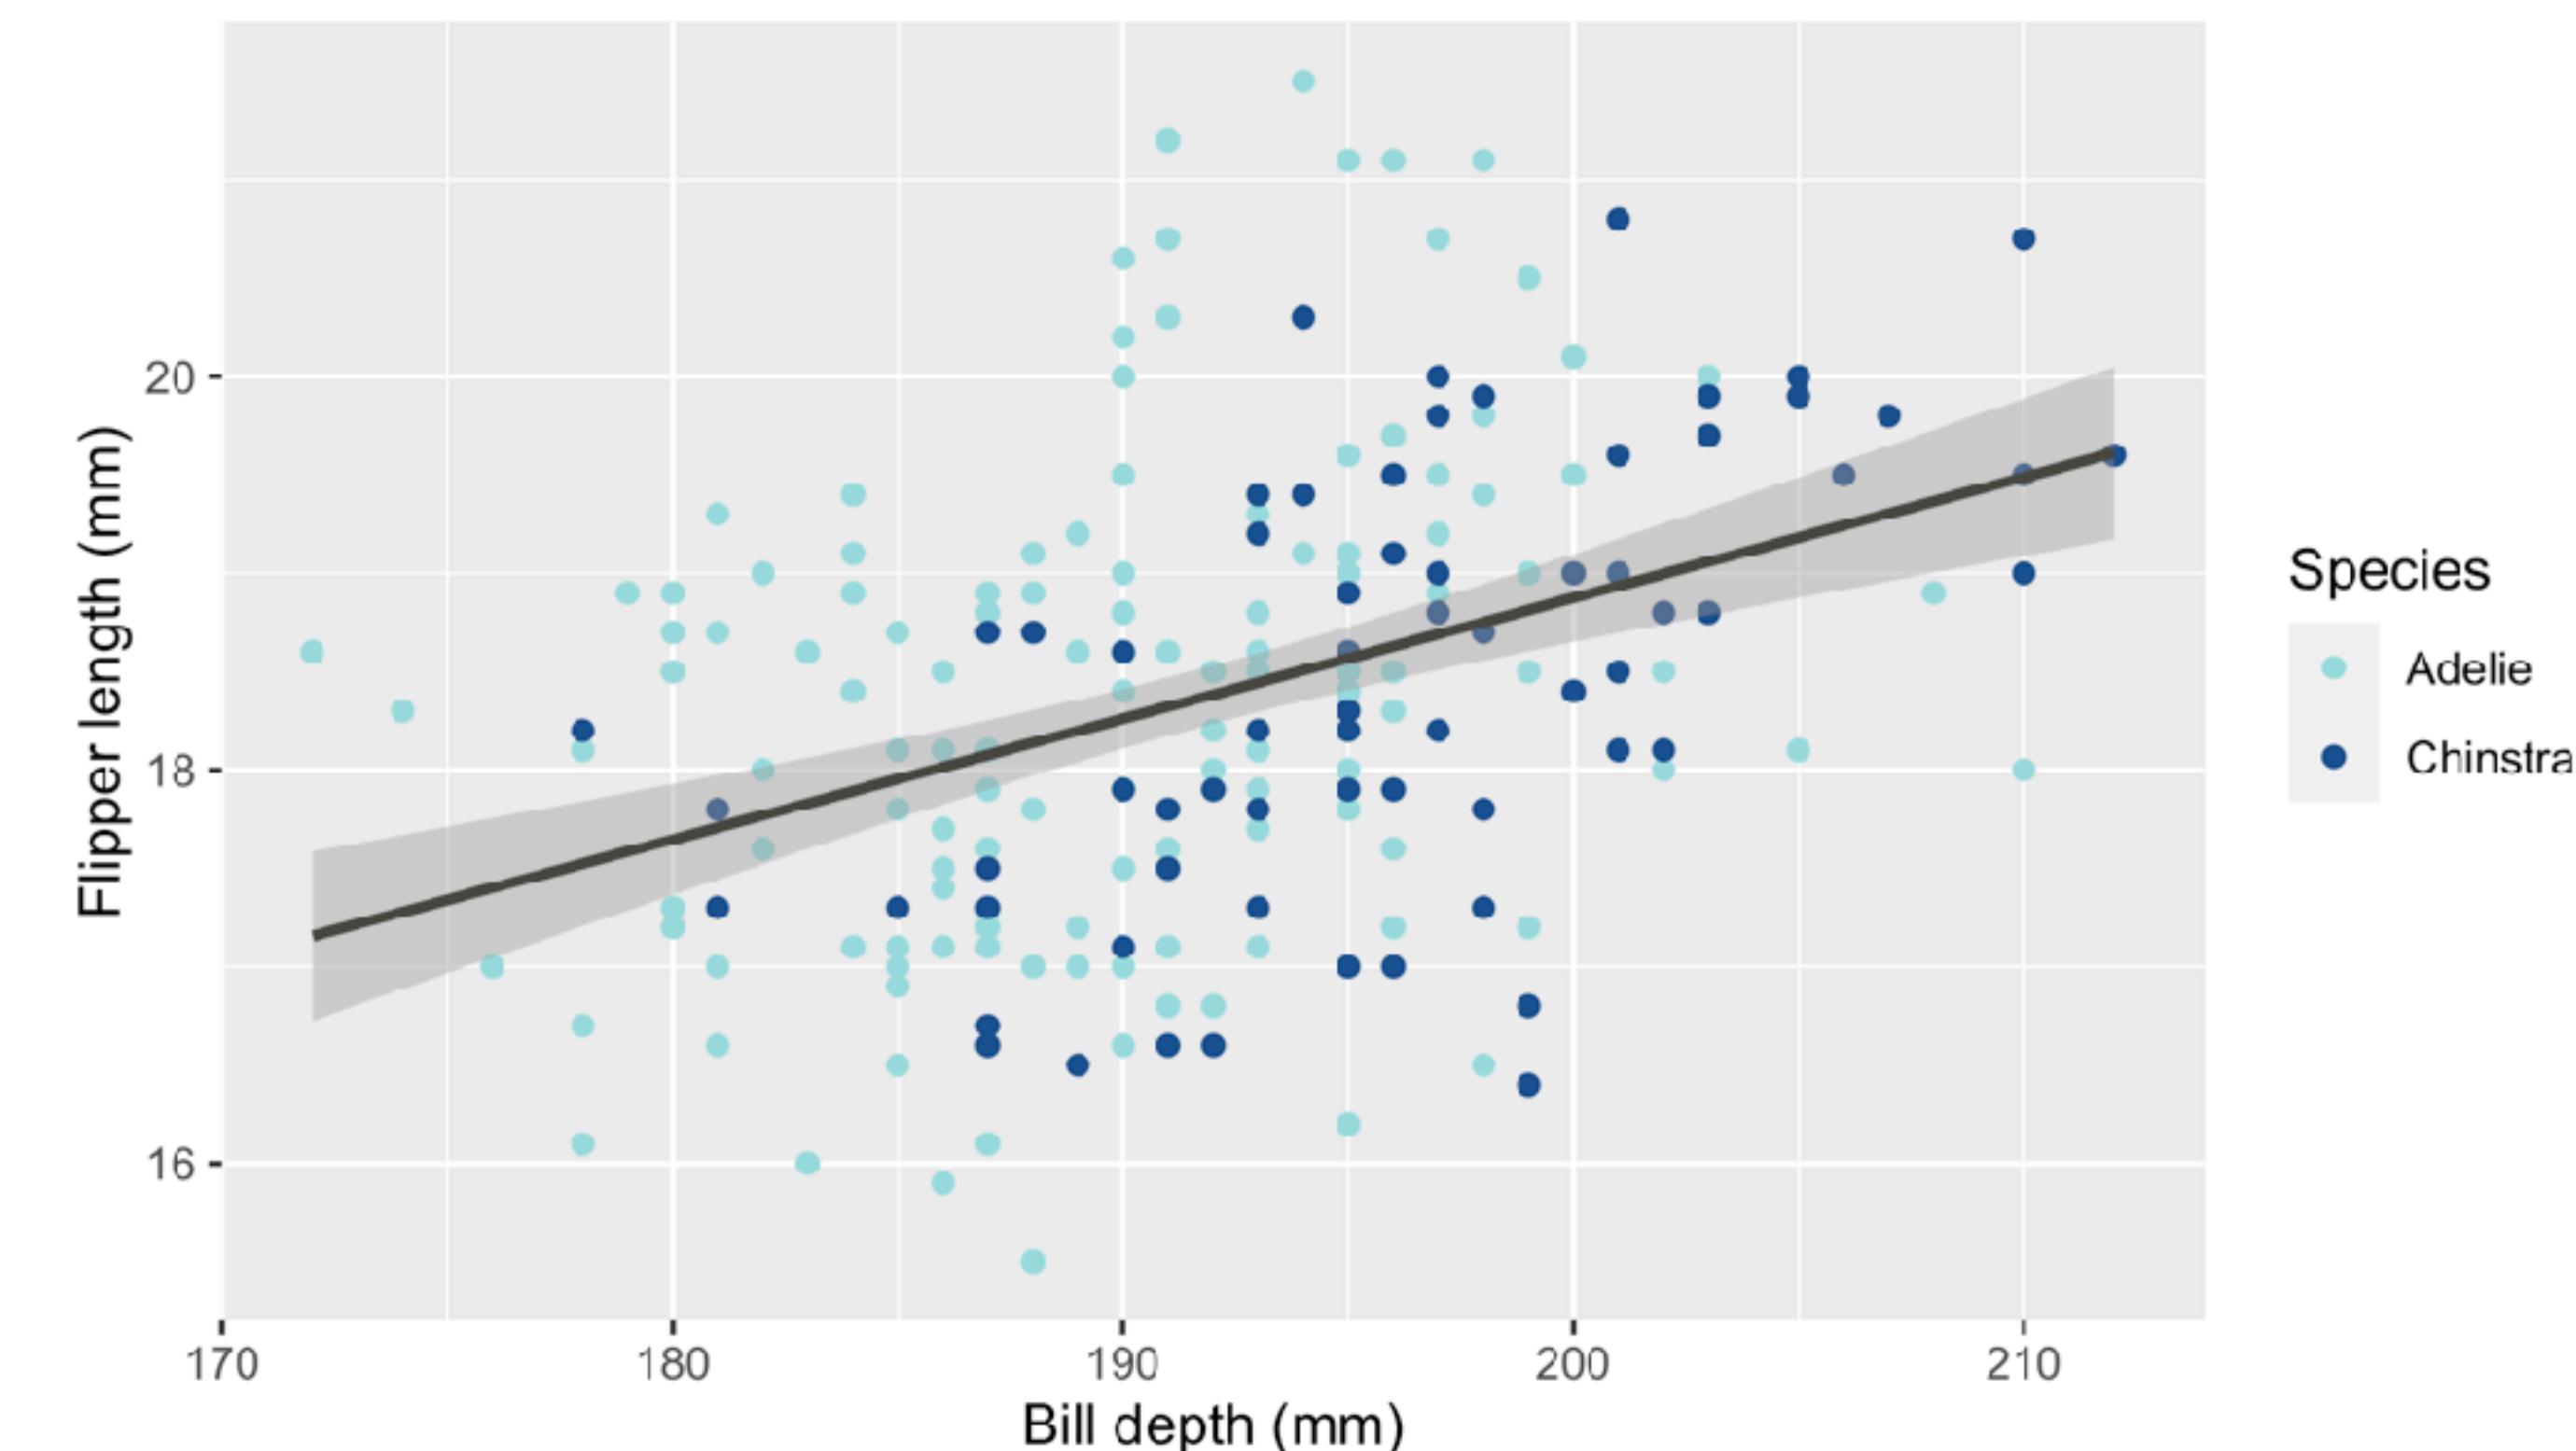


e.g.

Table 1. Regression output for predicting bill depth from flipper length.

term	estimate	std.error	statistic	p.value
(Intercept)	33.6	1.25	27.0	1.39e-86
flipper_length_mm	-0.0820	0.00618	-13.3	1.23e-32

Figure 2. Relationship between bill depth and flipper length.

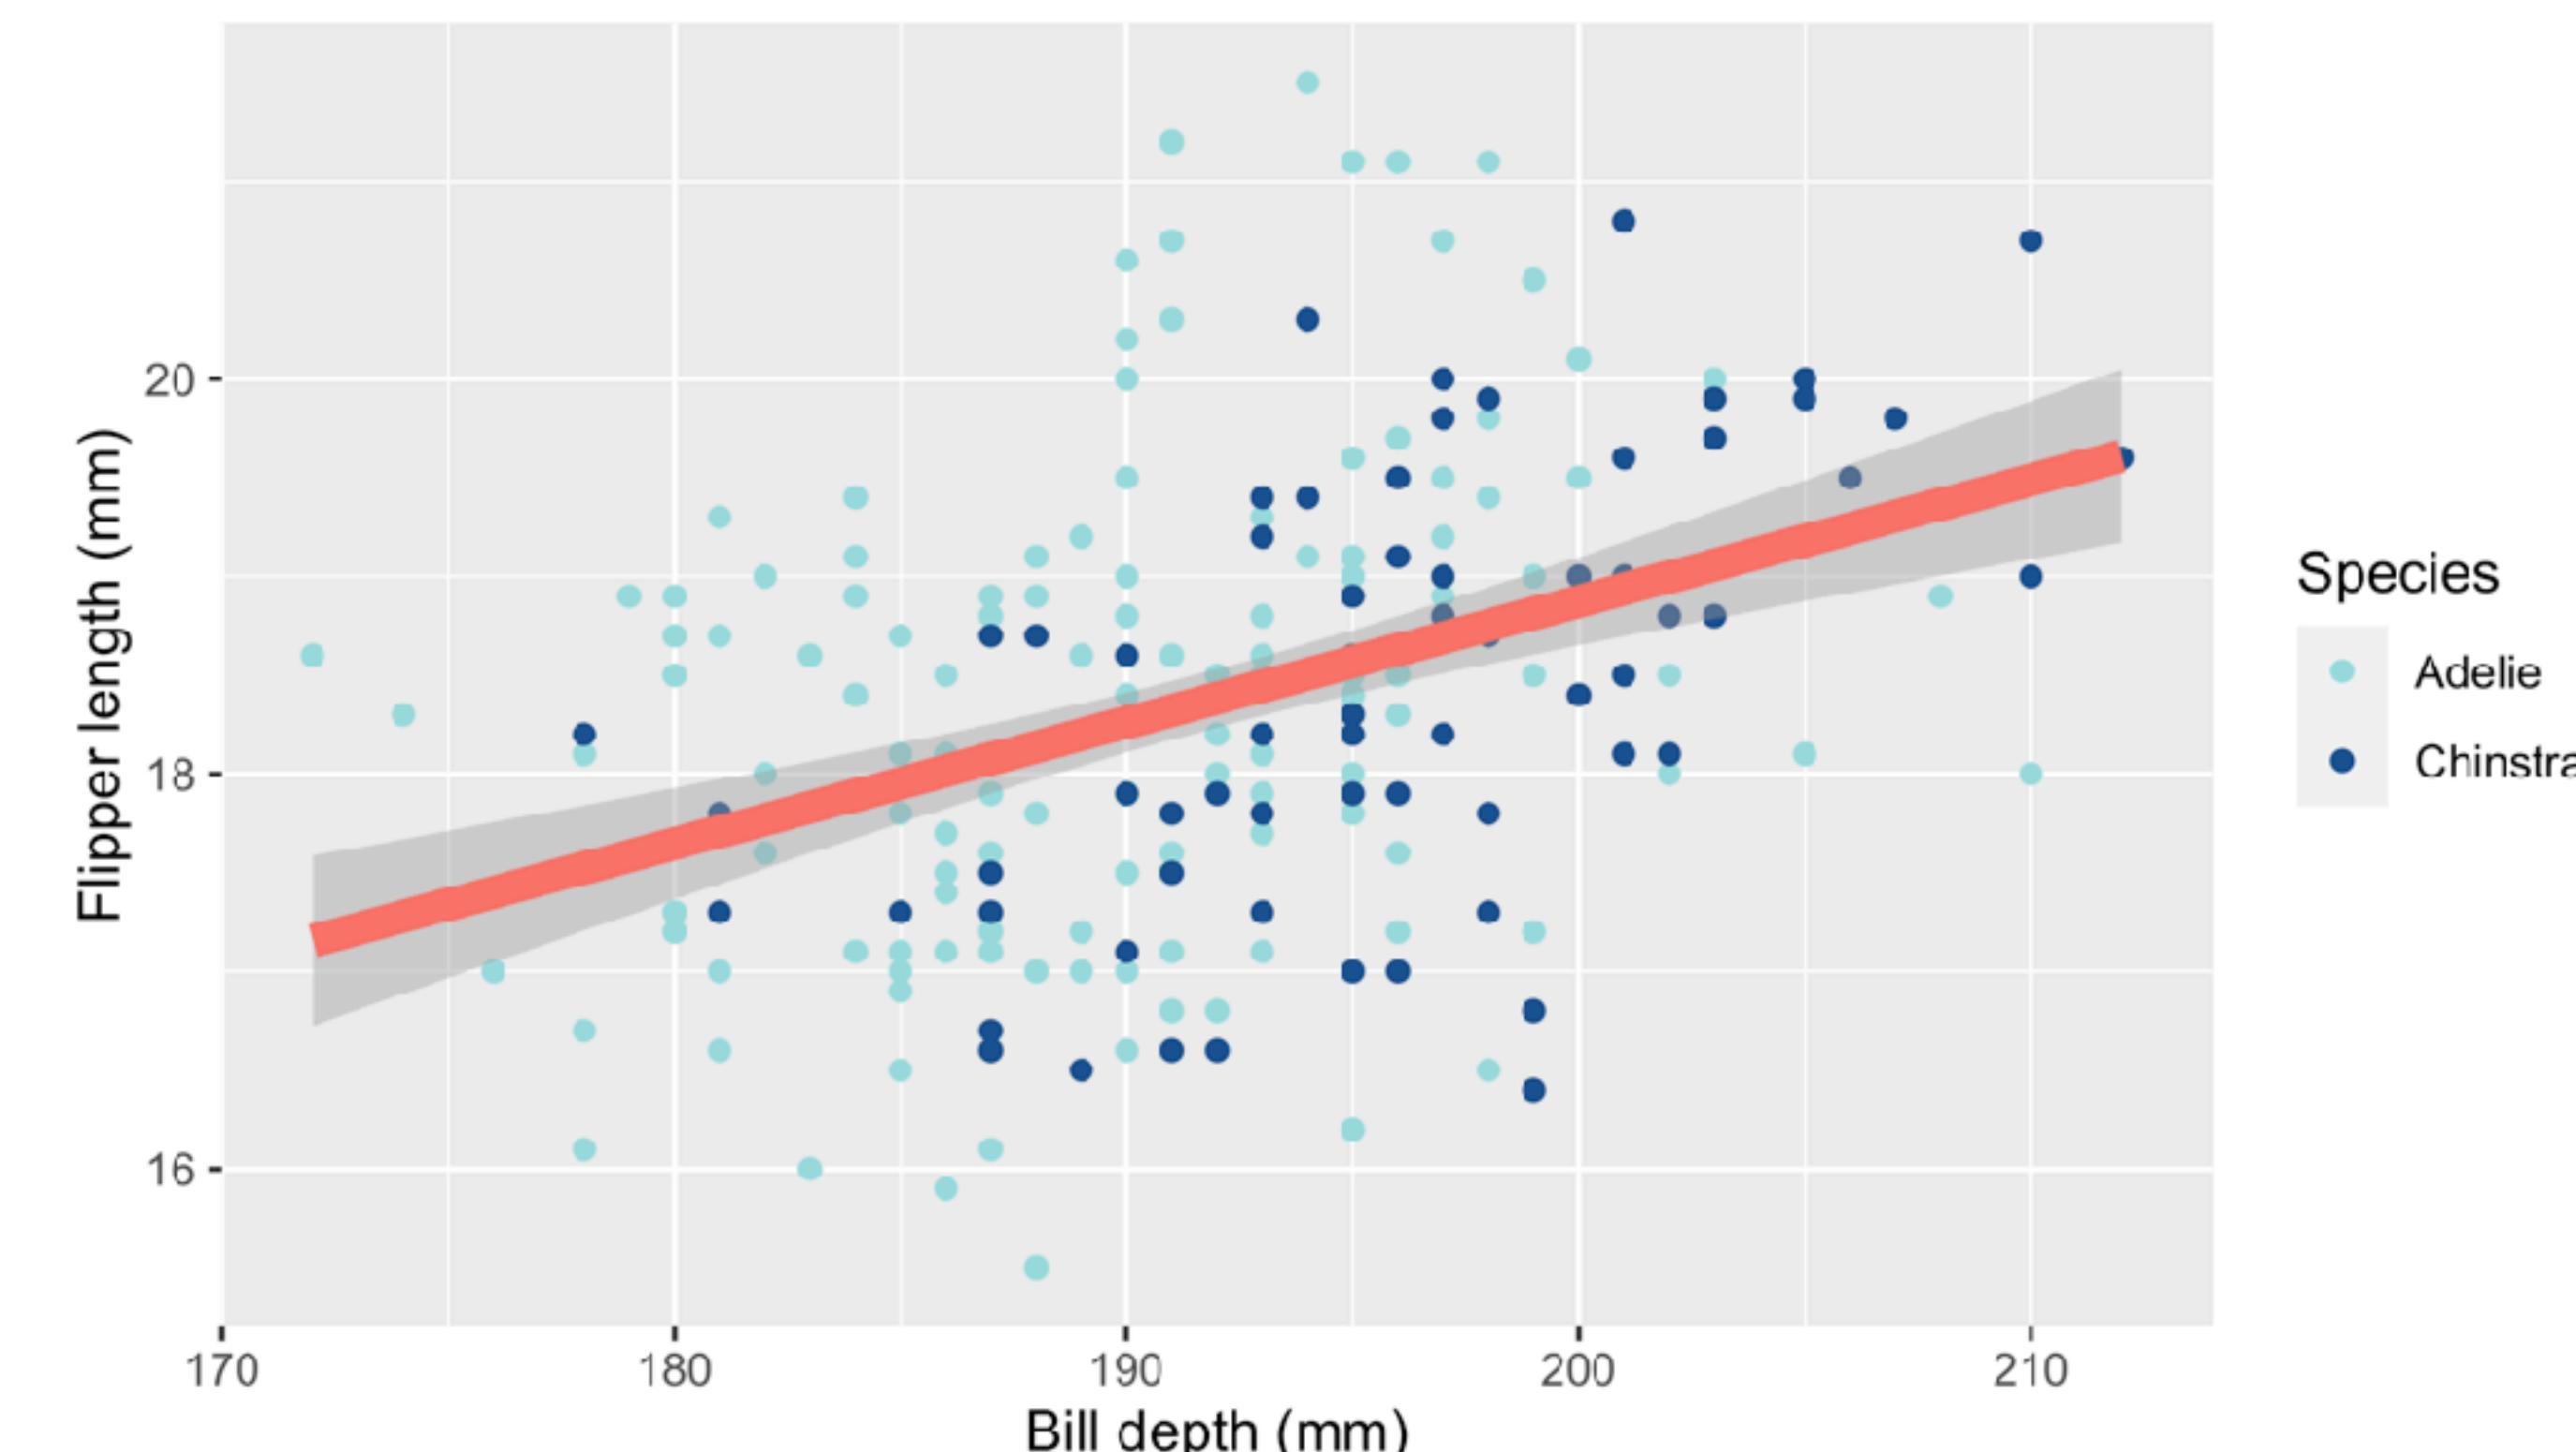


e.g.

Table 1. Regression output for predicting bill depth from flipper length.

term	estimate	std.error	statistic	p.value
(Intercept)	33.6	1.25	27.0	1.39e-86
flipper_length_mm	-0.0820	0.00618	-13.3	1.23e-32

Figure 2. Relationship between bill depth and flipper length.



analysis

report



```
# fit model  
model <- lm(bill_depth_mm ~ flipper_length_mm, data = penguins)  
  
# print model summary  
tidy(model)
```

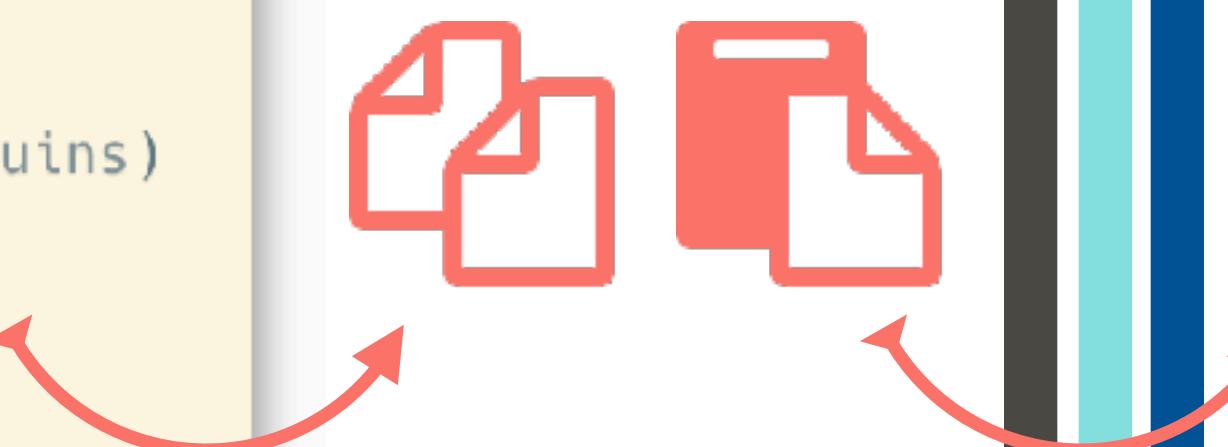


Table 1. Regression output for predicting bill depth from flipper length.

term	estimate	std.error	statistic	p.value
(Intercept)	33.6	1.25	27.0	1.39e-86
flipper_length_mm	-0.0820	0.00618	-13.3	1.23e-32

analysis

report



```
# visualize the relationship
ggplot(penguins) +
  geom_point(
    aes(x = bill_depth_mm, y = flipper_length_mm, color = species))
  +
  geom_smooth(
    aes(x = bill_depth_mm, y = flipper_length_mm),
    method = "lm")
```

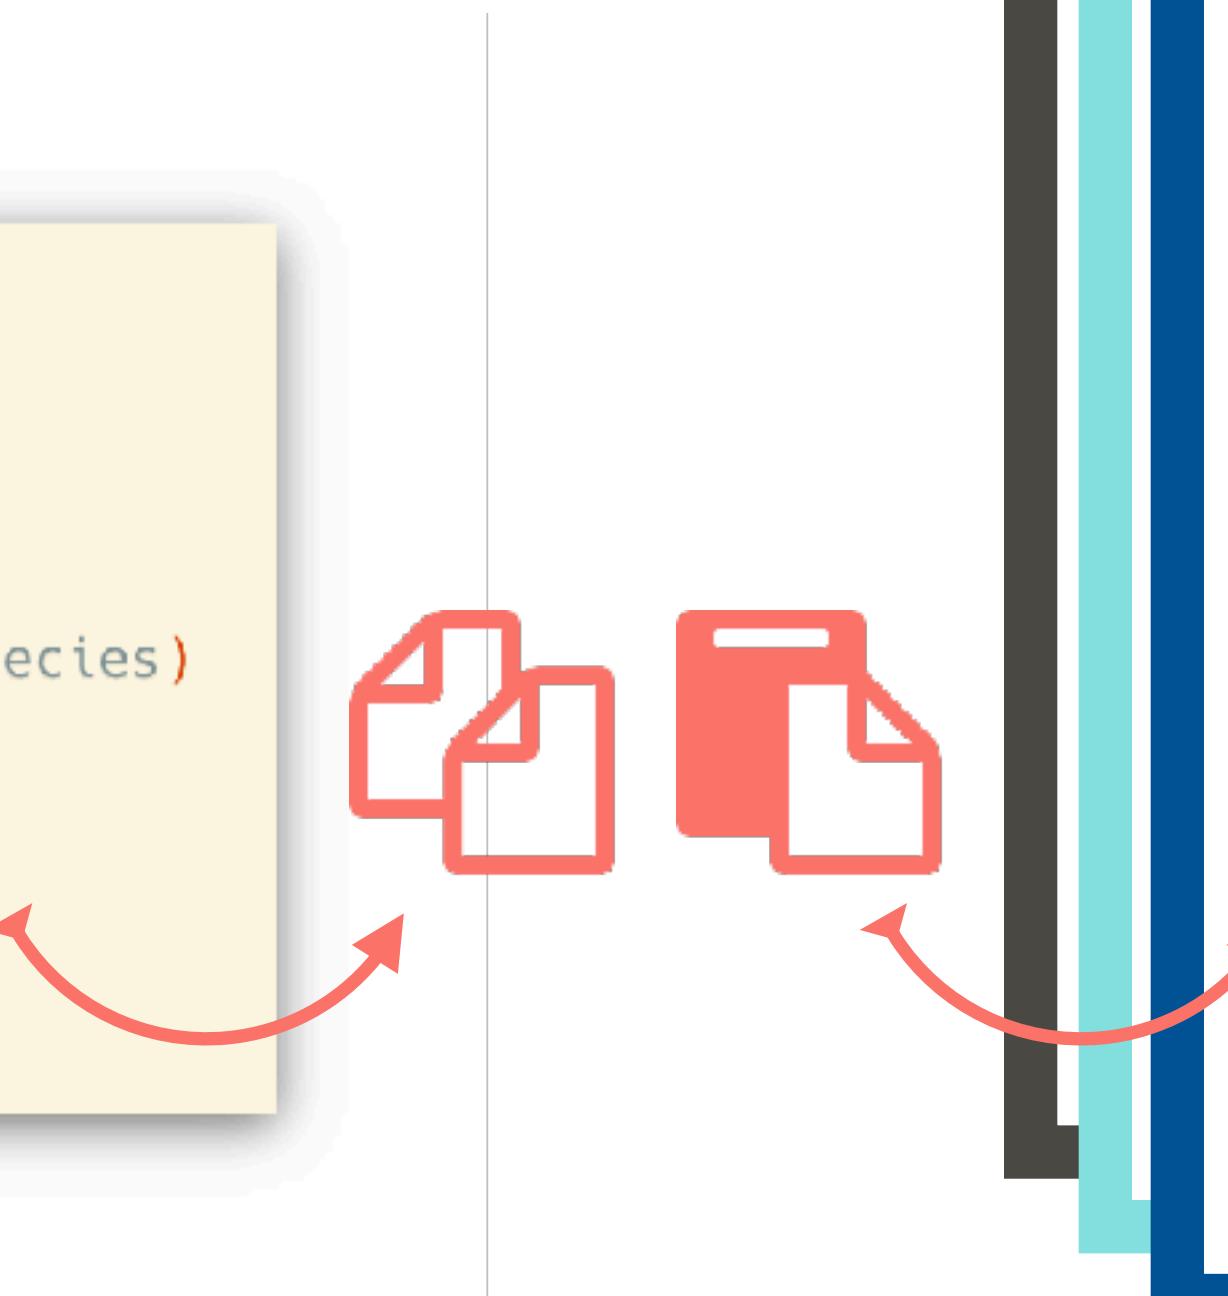
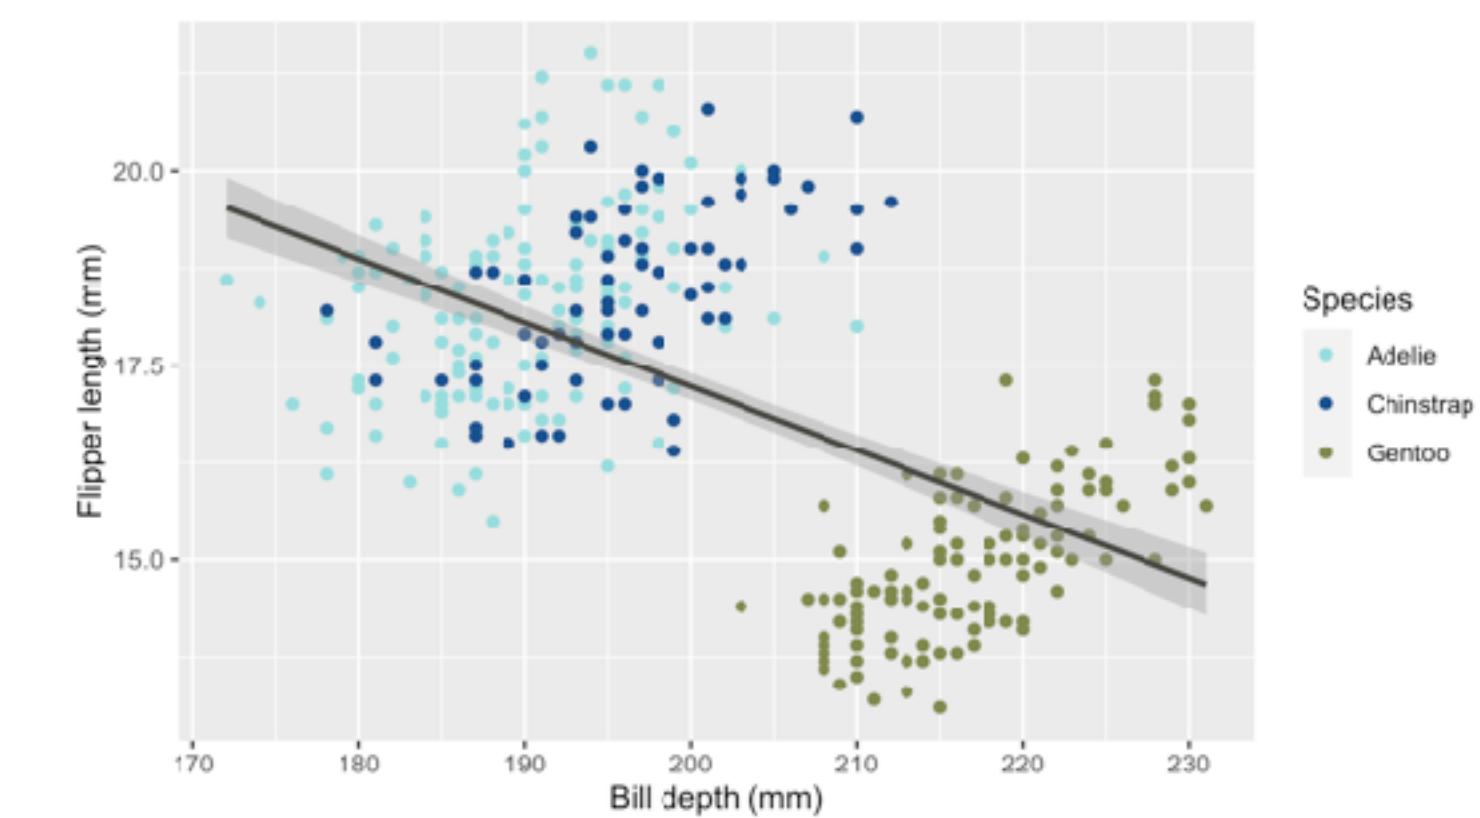


Table 1. Regression output for predicting bill depth from flipper length.

term	estimate	std.error	statistic	p.value
(Intercept)	33.6	1.25	27.0	1.39e-86
flipper_length_mm	-0.0820	0.00618	-13.3	1.23e-32

Figure 2. Relationship between bill depth and flipper length.



analysis

report



```
# filter out Gentoos
penguins_nongentoo <- penguins %>%
  filter(species != "Gentoo")

# visualize the relationship
ggplot(penguins_nongentoo) +
  geom_point(
    aes(x = bill_depth_mm, y = flipper_length_mm, color = species)
  ) +
  geom_smooth(
    aes(x = bill_depth_mm, y = flipper_length_mm),
    method = "lm"
  )
```

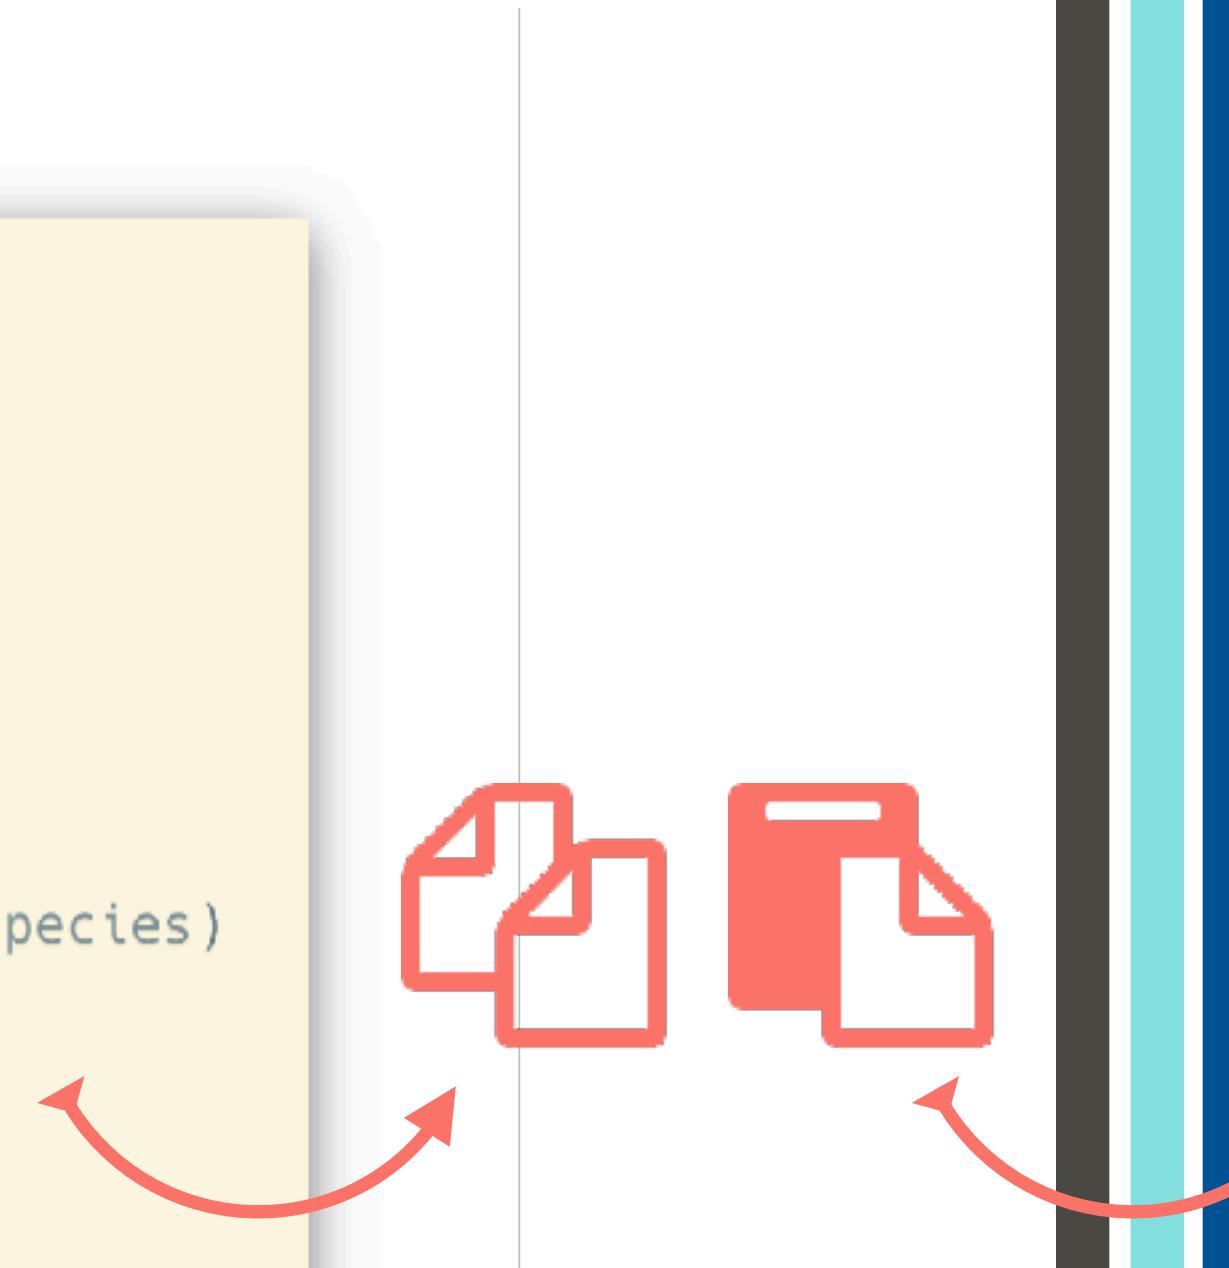


Table 1. Regression output for predicting bill depth from flipper length.

term	estimate	std.error	statistic	p.value
(Intercept)	33.6	1.25	27.0	1.39e-86
flipper_length_mm	-0.0820	0.00618	-13.3	1.23e-32

Figure 2. Relationship between bill depth and flipper length.

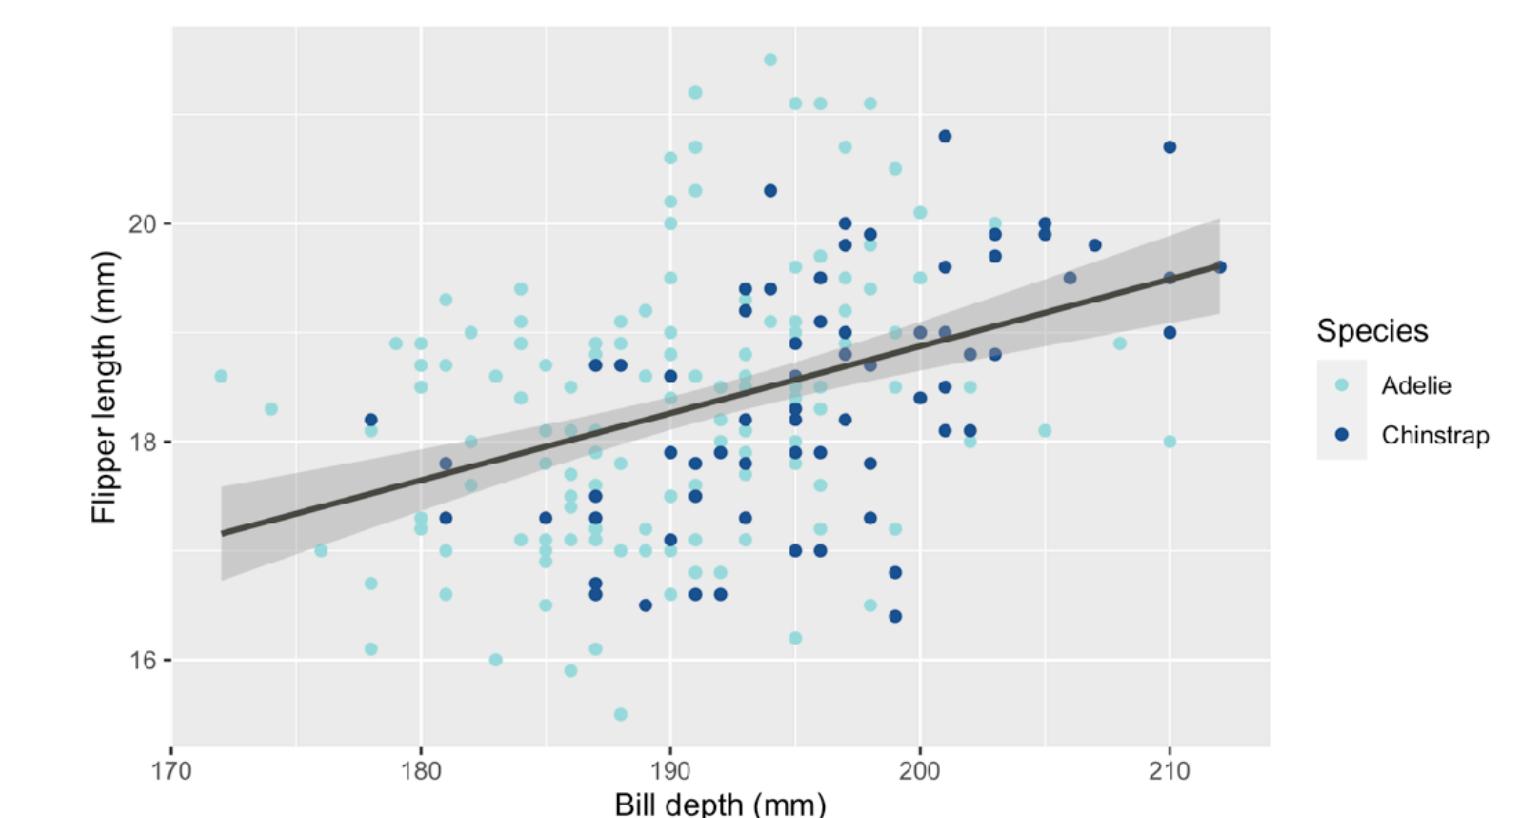
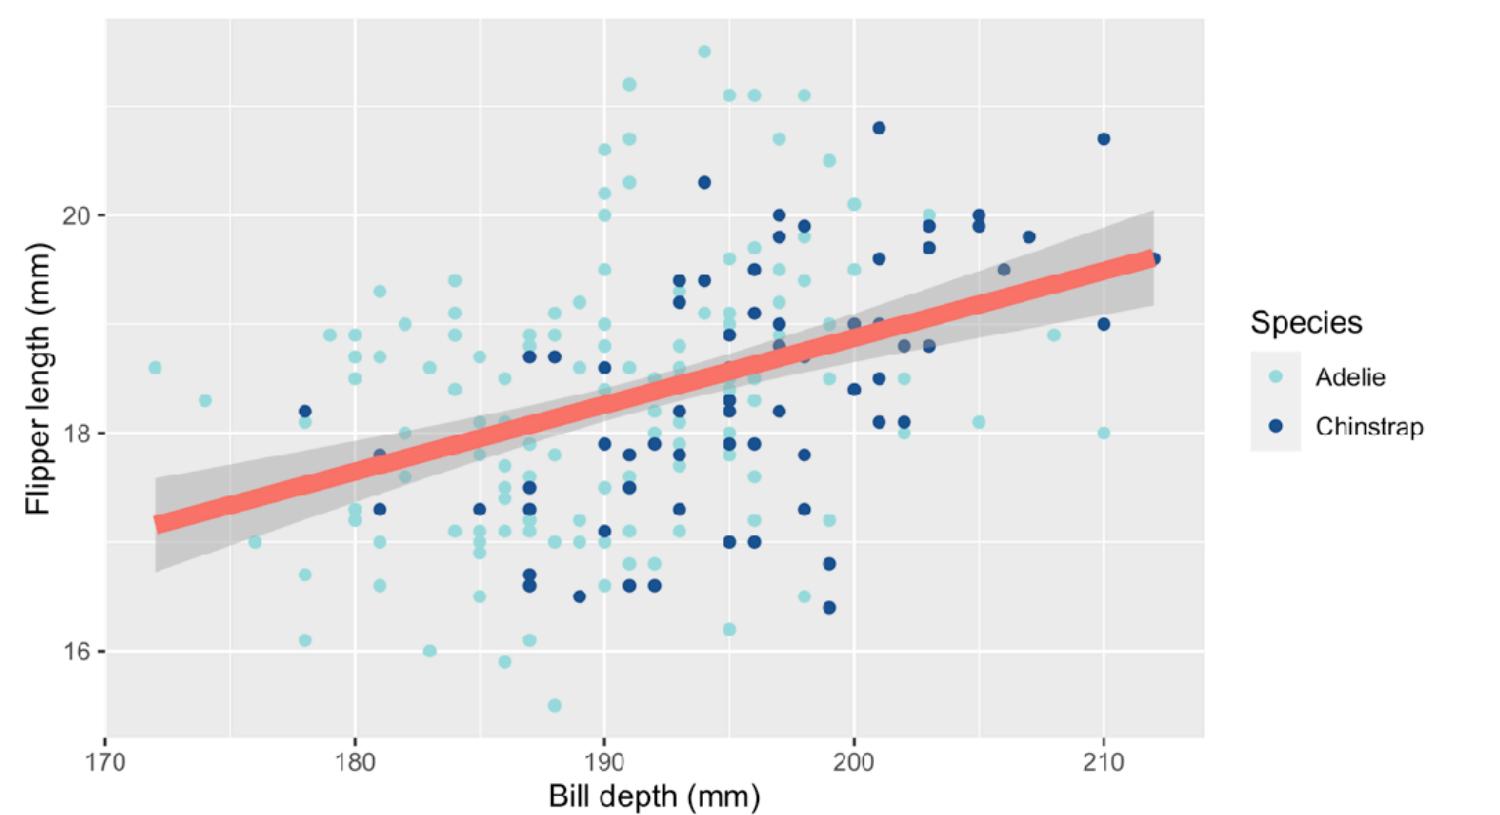
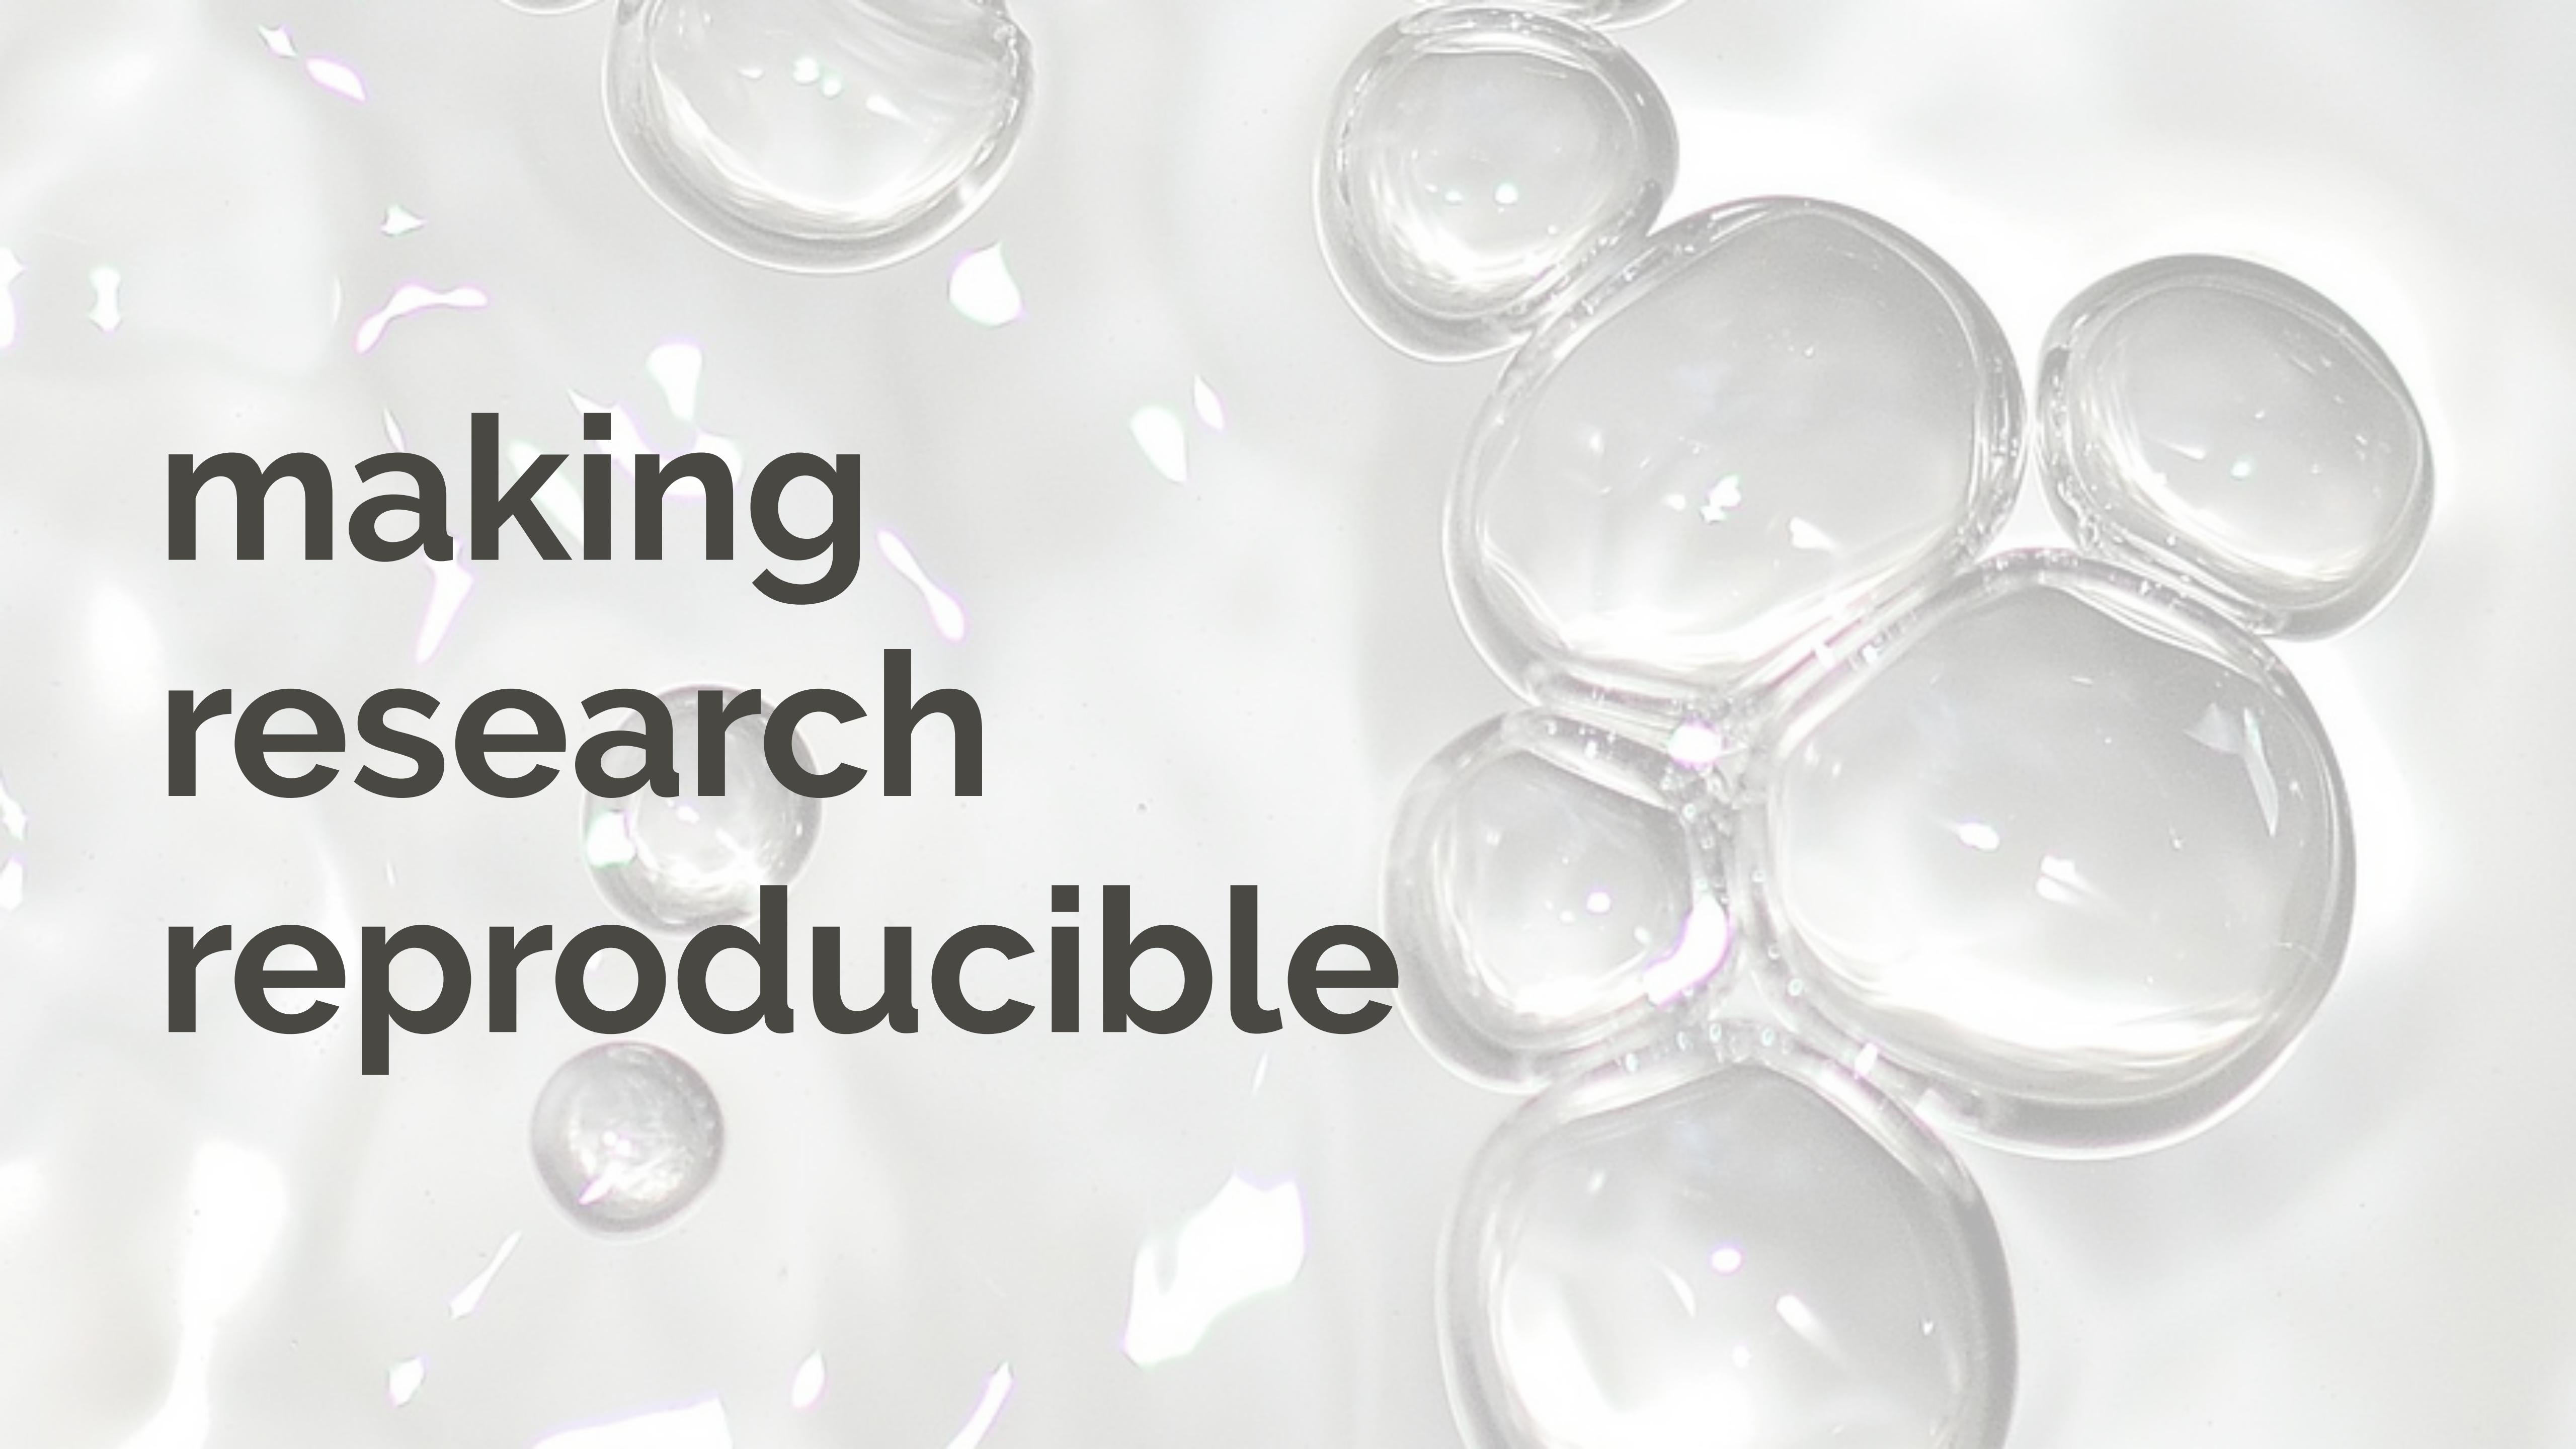


Table 1. Regression output for predicting bill depth from flipper length.

term	estimate	std.error	statistic	p.value
(Intercept)	33.6	1.25	27.0	1.39e-86
flipper_length_mm	-0.0820	0.00618	-13.3	1.23e-32

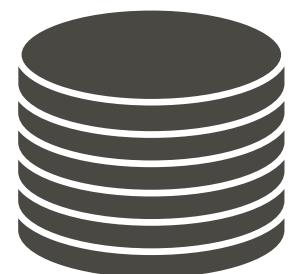
Figure 2. Relationship between bill depth and flipper length.





**making
research
reproducible**

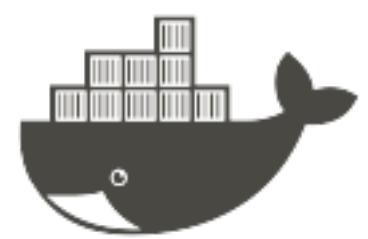
make



raw data



code & documentation to reproduce the analysis



specifications of your computational environment

available and accessible

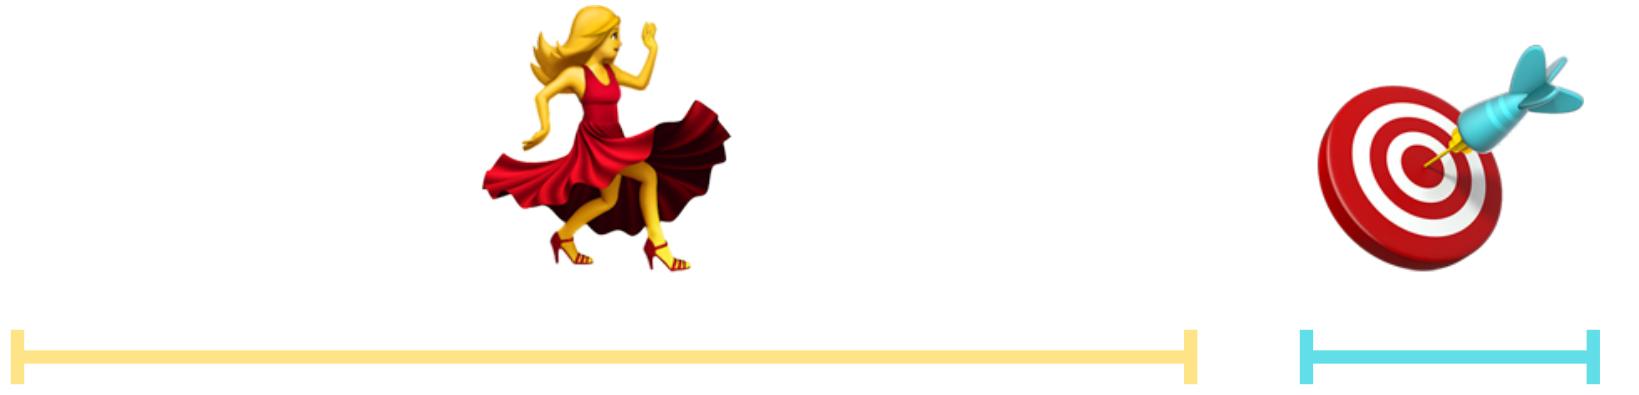
Peng, Roger. "The reproducibility crisis in science: A statistical counterattack." *Significance* 12.3 (2015): 30-32.

Gentleman, Robert, and Duncan Temple Lang. "Statistical analyses and reproducible research." *Journal of Computational and Graphical Statistics* 16.1 (2007): 1-23.

“The most important tool is
the **mindset**, when starting,
that the end product will be
reproducible.”

– Keith Baggerly

nobody,
not even yourself,
can recreate any part
of your analysis



push button
reproducibility
in published work

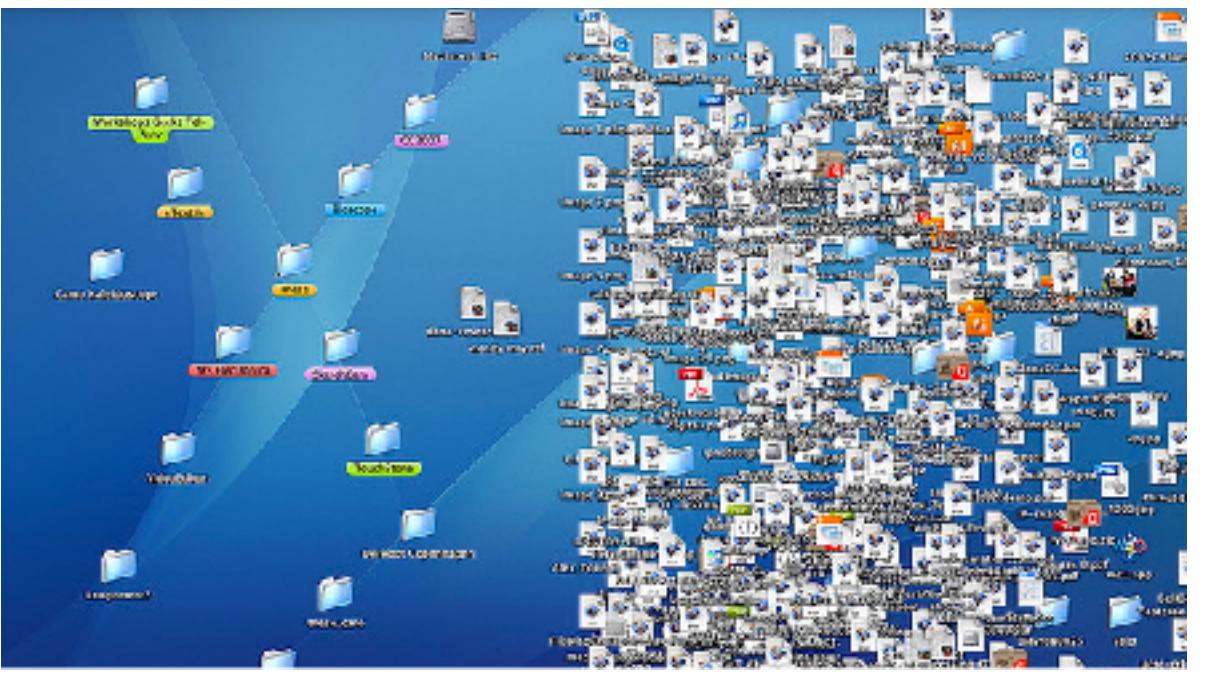
“There’s **no one-size-fits-all solution**
for computational reproducibility.”

8 principles
but the following[^] might help...

1

organize
your
project

level of organization



simpler analysis

-  raw-data
-  processed-data
-  manuscript
 - └ manuscript.Rmd

more complex analysis

-  raw-data
-  processed-data
-  scripts
-  figures
-  manuscript
 - └ manuscript.Rmd

stick with the conventions of your peers



write
READMEs
liberally



raw-data

- └ README.md
- └ airlines.csv
- └ airports.csv
- └ flights.csv
- └ planes.csv
- └ weather.csv



processed-data



scripts



figures



manuscript



README

This folder contains the raw data for the project.

All datasets were downloaded from openflights.org/data.html on 2019-04-01.

- airlines: Airline names
- airports: Airports metadata
- flights: Flight data
- planes: Plane metadata
- weather: Hourly weather data



keep data
tidy &
machine readable

Student	Exam Grade		
Name	1	2	Major
Barney Donaldson	89	76	Data Science, Public Policy
Clay Whelan	67	83	Public Policy
Simran Bass	82	90	Statistics
Chante Munro	45	72	Political Science, Statistics
Gabrielle Cherry	32	79	.
Kush Piper	98	sick	Statistics
Faizan Ratliff	82	75	Data Science
Torin Ruiz	70	80	Sociology, Statistics
Reiss Richardson	missed exam	34	Neuroscience
Ajwa Cochran	50	65	Data Science

→ record
code +
document
non-code
steps +
write
tests

name	exam_1	exam_2	first_major	second_major	participation
Barney Donaldson	89	76	Data Science	Public Policy	ok
Clay Whelan	67	83	Public Policy	NA	ok
Simran Bass	82	90	Statistics	NA	ok
Chante Munro	45	72	Political Science	Statistics	low
Gabrielle Cherry	32	79	NA	NA	ok
Kush Piper	98	NA	Statistics	NA	ok
Faizan Ratliff	82	75	Data Science	NA	ok
Torin Ruiz	70	80	Sociology	Statistics	ok
Reiss Richardson	NA	34	Neuroscience	NA	low
Ajwa Cochran	50	65	Data Science	NA	low

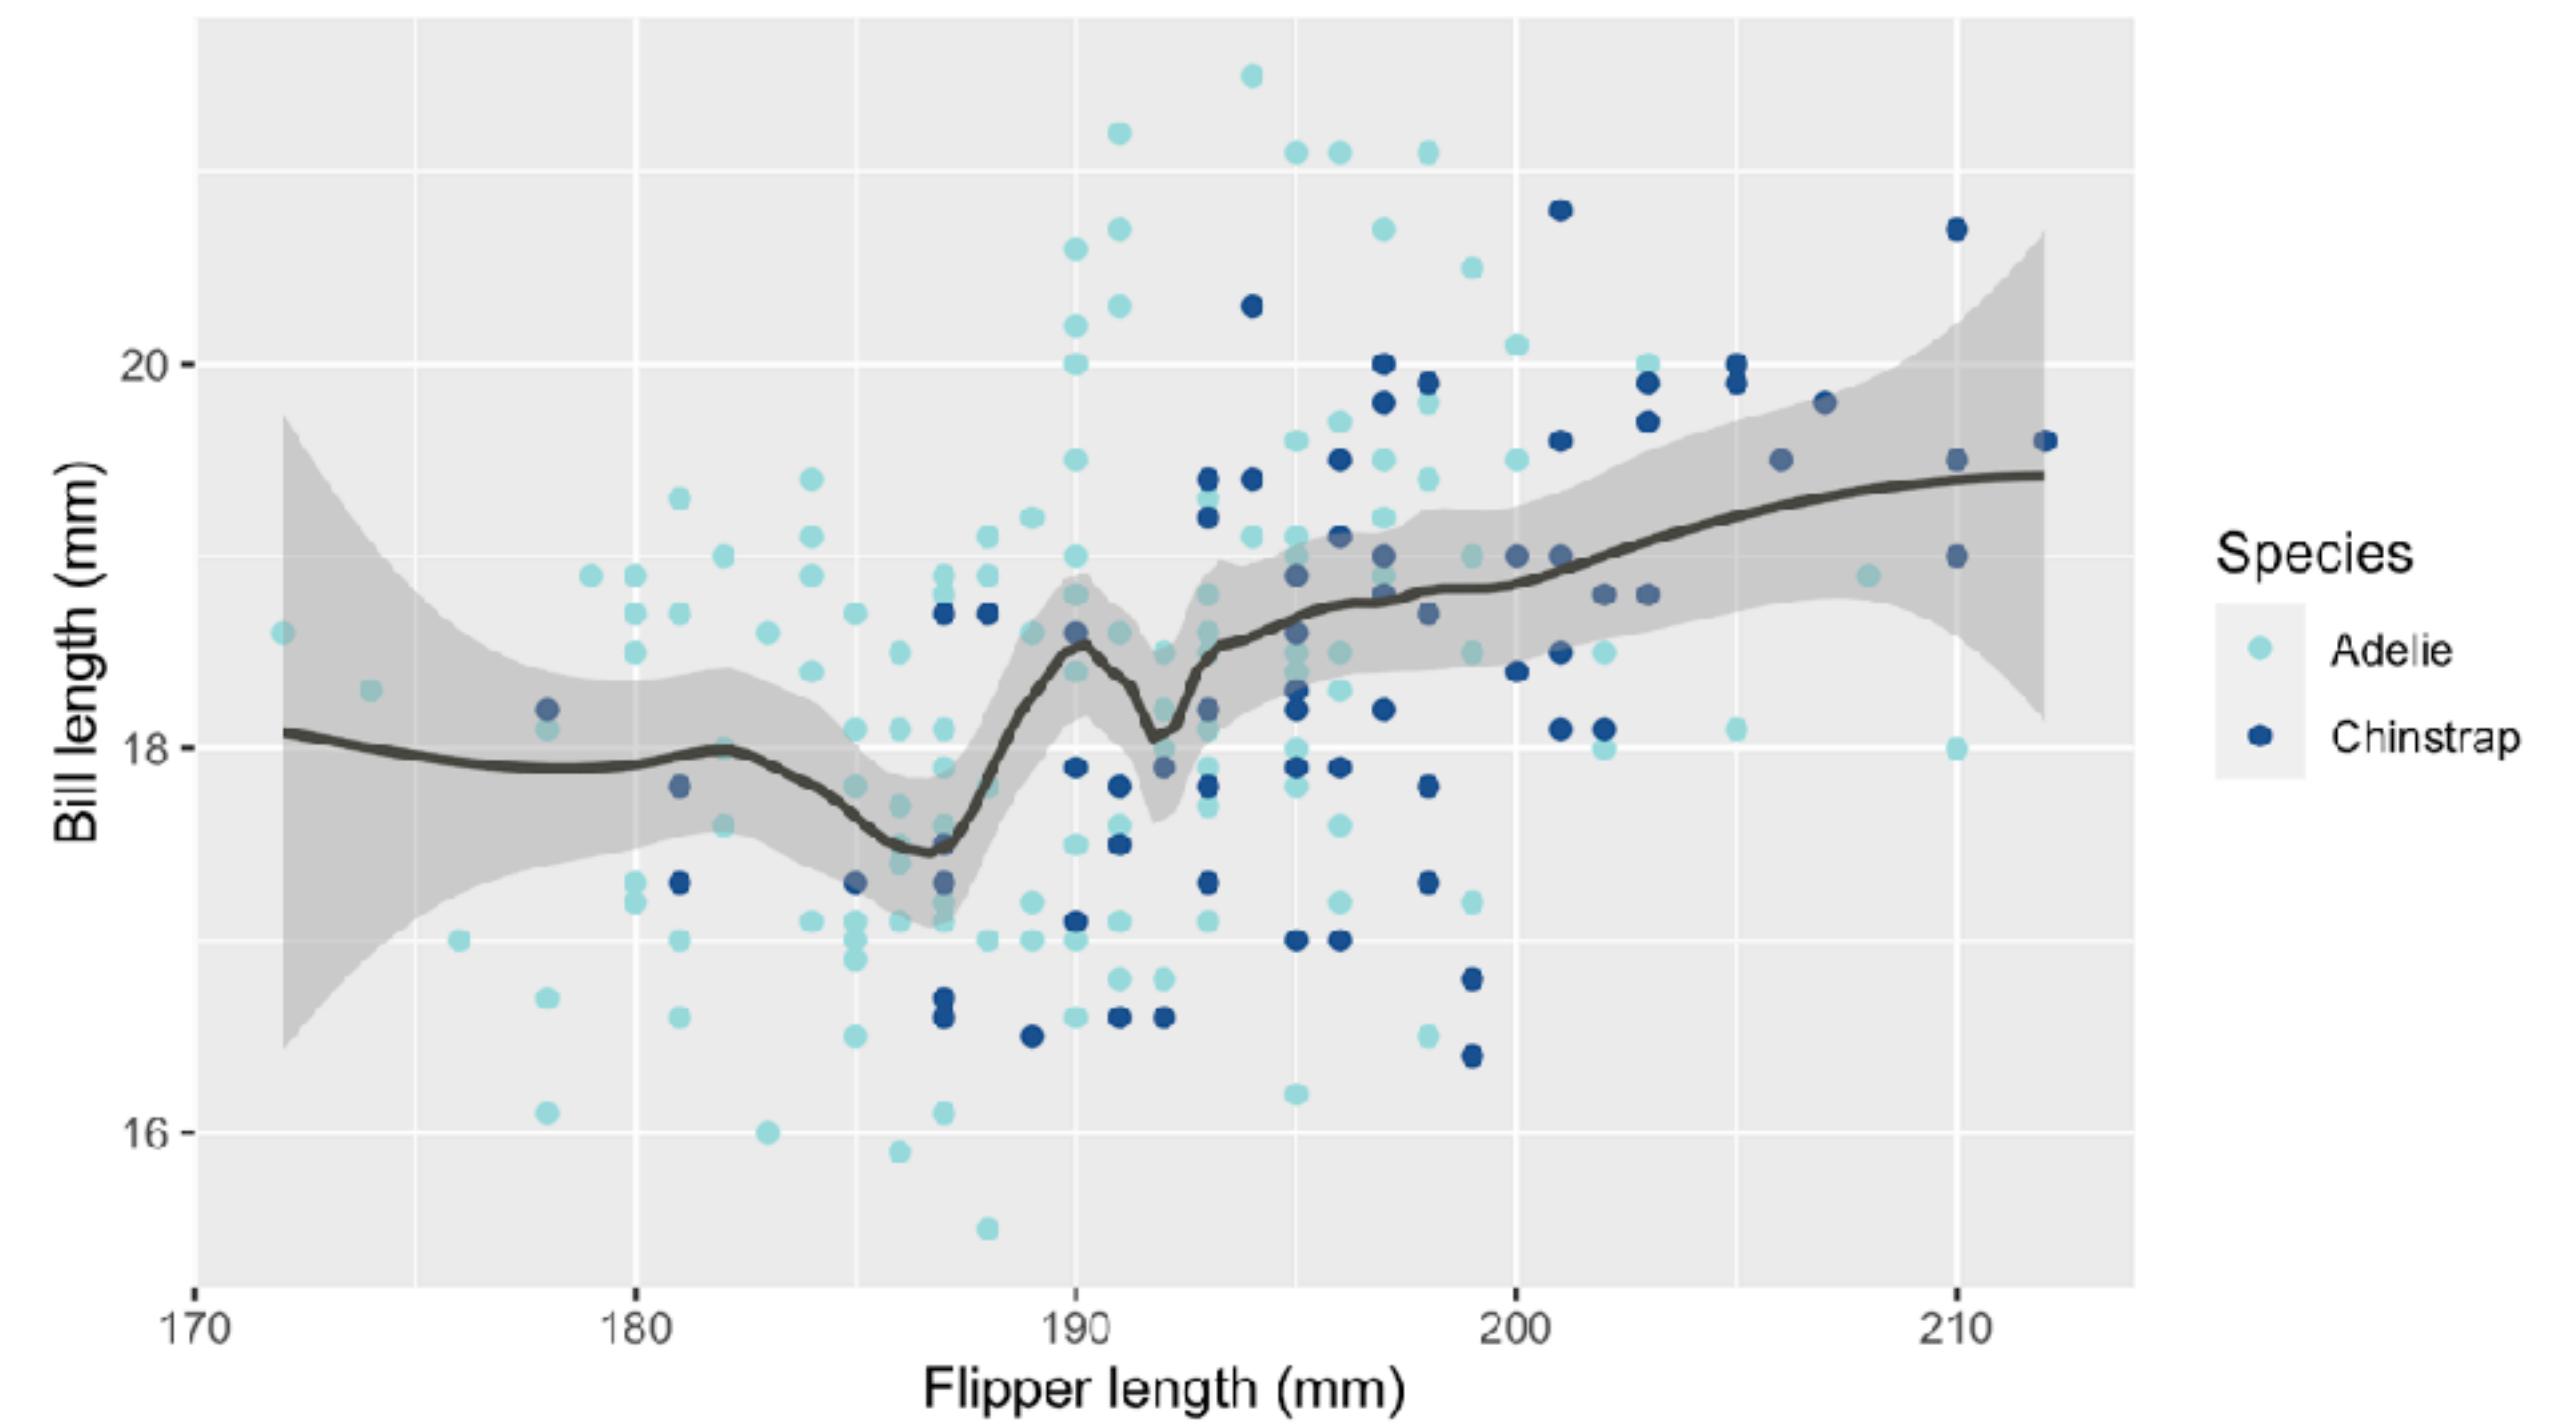
Low participation

1

comment
your
code



```
# use loess smoothing
ggplot(penguins_nongentoo) +
  geom_point(
    aes(x = flipper_length_mm, y = bill_depth_mm, color = species)
  ) +
  geom_smooth(
    aes(x = flipper_length_mm, y = bill_depth_mm),
    method = "loess", span = 0.375
  )
```





use
literate
programming

tab1-fig2.Rmd

Normal | B I U </> | X | : 1= , | @ | Format | Insert | Table

```
---
```

```
title: "Table 1 matches Figure 2!"  
author: "Mine Çetinkaya-Rundel"  
date: "`r Sys.Date()`"  
output:  
  bookdown::html_document2:  
    fig_caption: yes  
bibliography: references.bib  
---
```

```
{r chunk-options, include=FALSE}  
knitr::opts_chunk$set(echo = FALSE, message = FALSE)
```

```
{r setup}  
library(broom)          # for tidy model output  
library(glue)           # for gluing strings to data  
library(knitr)           # for kable  
library(palmerpenguins)  # for data  
library(tidyverse)        # for data wrangling and visualisation
```

In this report we evaluate the relationship between relationship between bill depth and flipper length of penguins. The data come from [@gorman2014](#).

The diagram shows a side view of a penguin's head. A horizontal double-headed arrow labeled 'Bill length' spans the width of the bill from the tip to the base. A vertical double-headed arrow labeled 'Bill depth' measures the thickness of the bill at its widest point.

Exploratory data analysis | R Markdown

improve-repro-workflow-reproducibilitea-2020 - RStudio

tab1-fig2.Rmd x Addins

Environment History Connections Tutorial

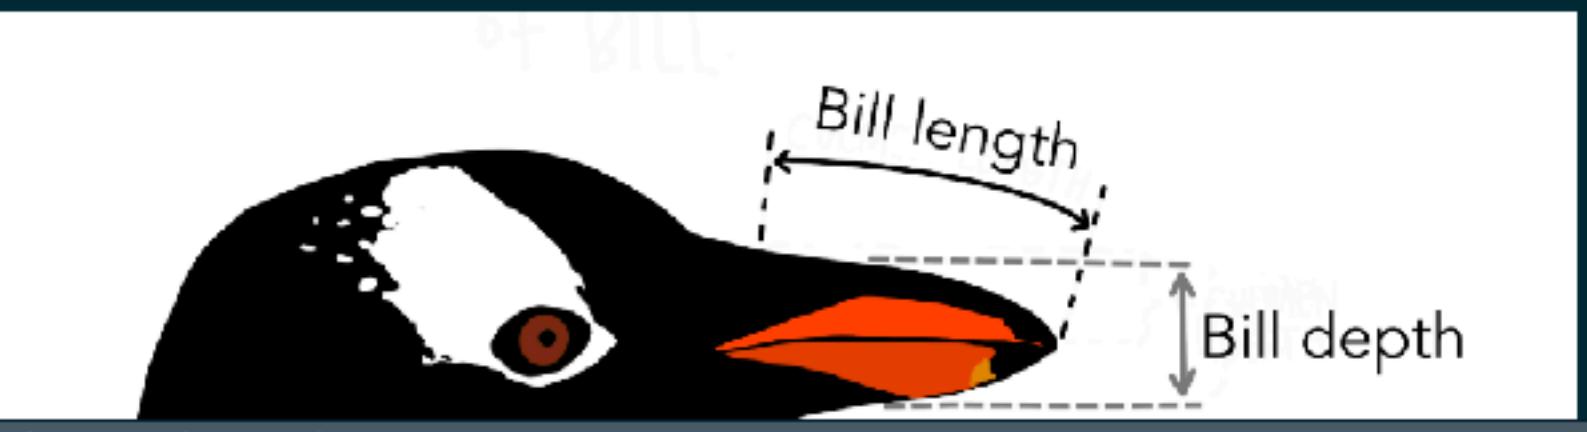
Files Plots Packages Help Viewer

Publish

Normal B I U Insert Table

```
---  
title: "Table 1 matches Figure 2!"  
author: "Mine Çetinkaya-Rundel"  
date: "r Sys.Date()"  
output:  
  bookdown::html_document2:  
    fig_caption: yes  
bibliography: references.bib  
---  
  
{r chunk-options, include=FALSE}  
knitr::opts_chunk$set(echo = FALSE, message = FALSE)  
  
{r setup}  
library(broom)      # for tidy model output  
library(glue)        # for gluing strings to data  
library(knitr)       # for kable  
library(palmerpenguins) # for data  
library(tidyverse)    # for data wrangling and visualisation
```

In this report we evaluate the relationship between bill depth and flipper length of penguins. The data come from @gorman2014.



Exploratory data analysis R Markdown

Console

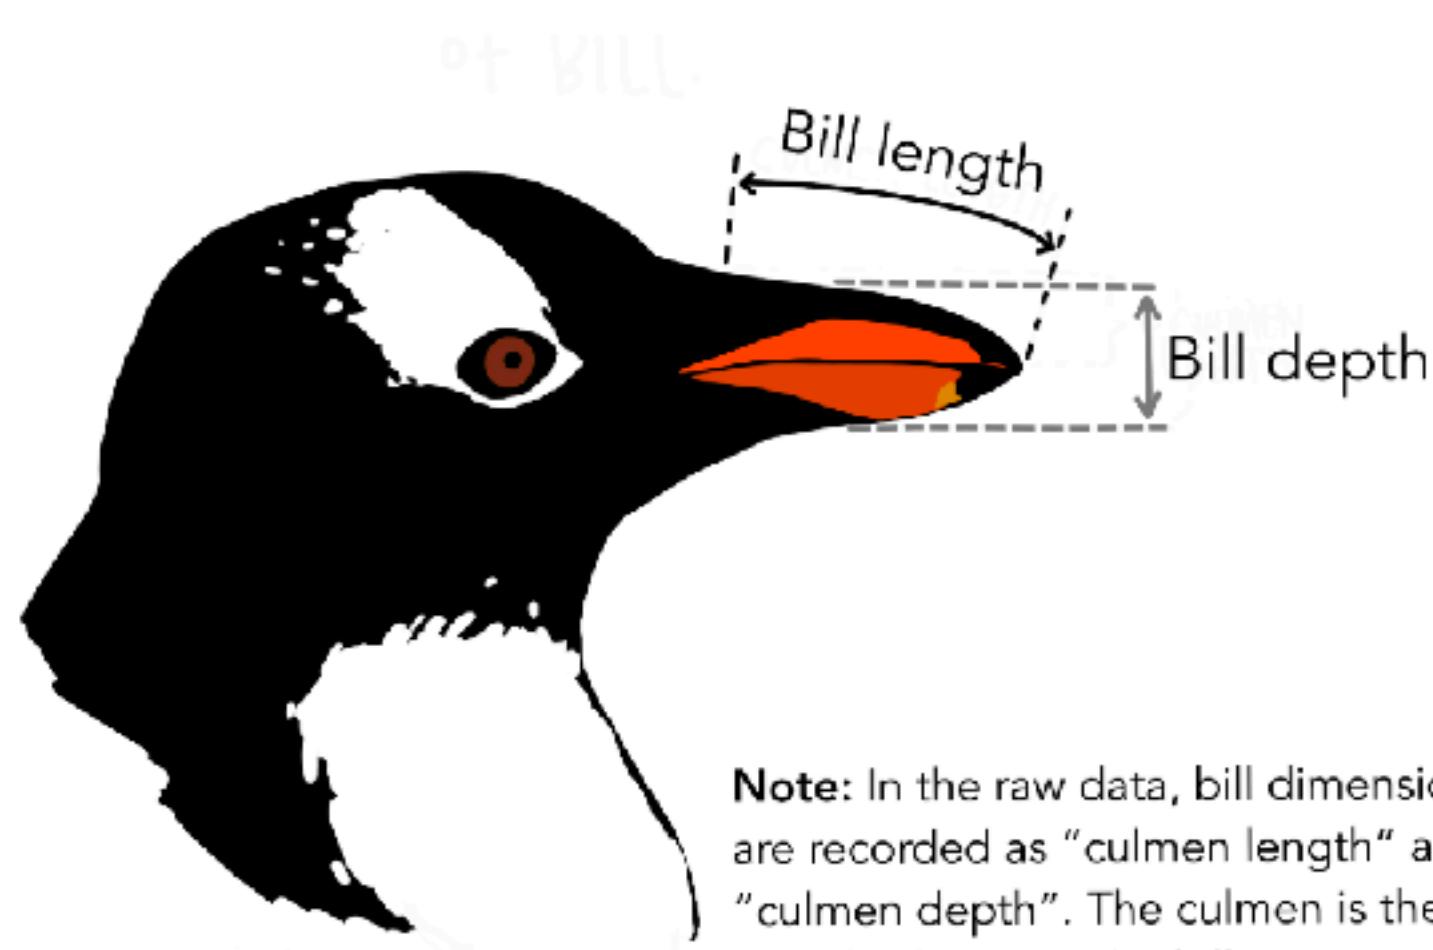
Table 1.1 shows some summary statistics.

Table 1.1: Summary statistics

Table 1 matches Figure 2!

Mine Çetinkaya-Rundel
2020-11-10

In this report we evaluate the relationship between bill depth and flipper length of penguins. The data come from Gorman, Williams, and Fraser (2014).



Note: In the raw data, bill dimensions are recorded as "culmen length" and "culmen depth". The culmen is the dorsal ridge atop the bill.

The original dataset has 3 species (Adelie, Gentoo, and Chinstrap), but we will only work with Adelie and Chinstrap species.

1 Exploratory data analysis

demo

rmarkdown

more resources . . .

- ▶ Learn more about **R Markdown**:
 - ▶ Documentation: rmarkdown.rstudio.com
 - ▶ Book: bookdown.org/yihui/rmarkdown
 - ▶ Book: bookdown.org/yihui/rmarkdown-cookbook
- ▶ Learn more about the **visual editor**:
 - ▶ Documentation: rstudio.github.io/visual-markdown-editing



use
version
control

The screenshot shows the RStudio interface with an R Markdown document open. The code chunk contains YAML front matter and R code for setting up packages and options. A figure of a penguin with measurement annotations is included.

```
---
```

```
title: "Table 1 matches Figure 2!"  
author: "Mine Çetinkaya-Rundel"  
date: "^r Sys.Date()"  
output:  
  bookdown::html_document2:  
    fig_caption: yes  
bibliography: references.bib
```

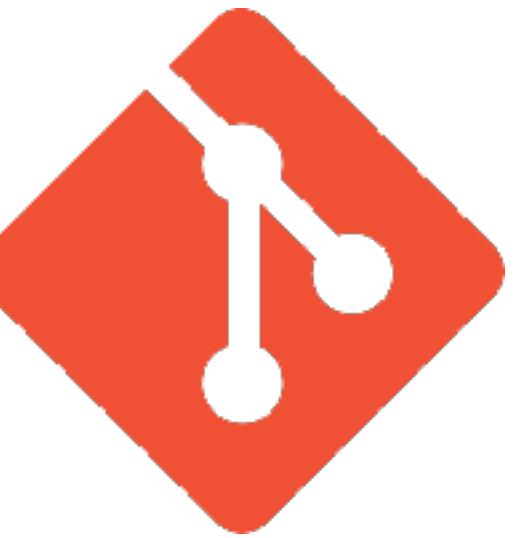
```
{r chunk-options, include=FALSE}  
knitr::opts_chunk$set(echo = FALSE, message = FALSE)
```

```
{r setup}  
library(broom)      # for tidy model output  
library(glue)       # for gluing strings to data  
library(knitr)      # for kable  
library(palmerpenguins) # for data  
library(tidyverse)   # for data wrangling and visualisation
```

In this report we evaluate the relationship between relationship between bill depth and flipper length of penguins. The data come from [@gorman2014](#).

A small image of a penguin's head and upper body. Two dashed arrows point from text labels to specific parts of the penguin's beak: one arrow points horizontally to the right and is labeled 'Bill length', and another arrow points vertically downwards and is labeled 'Bill depth'.

changes
tracked by



hosted
on

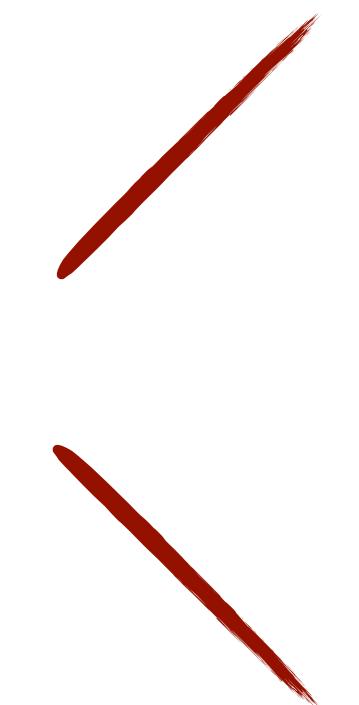


2

Git workflows

GitHub first

Local first

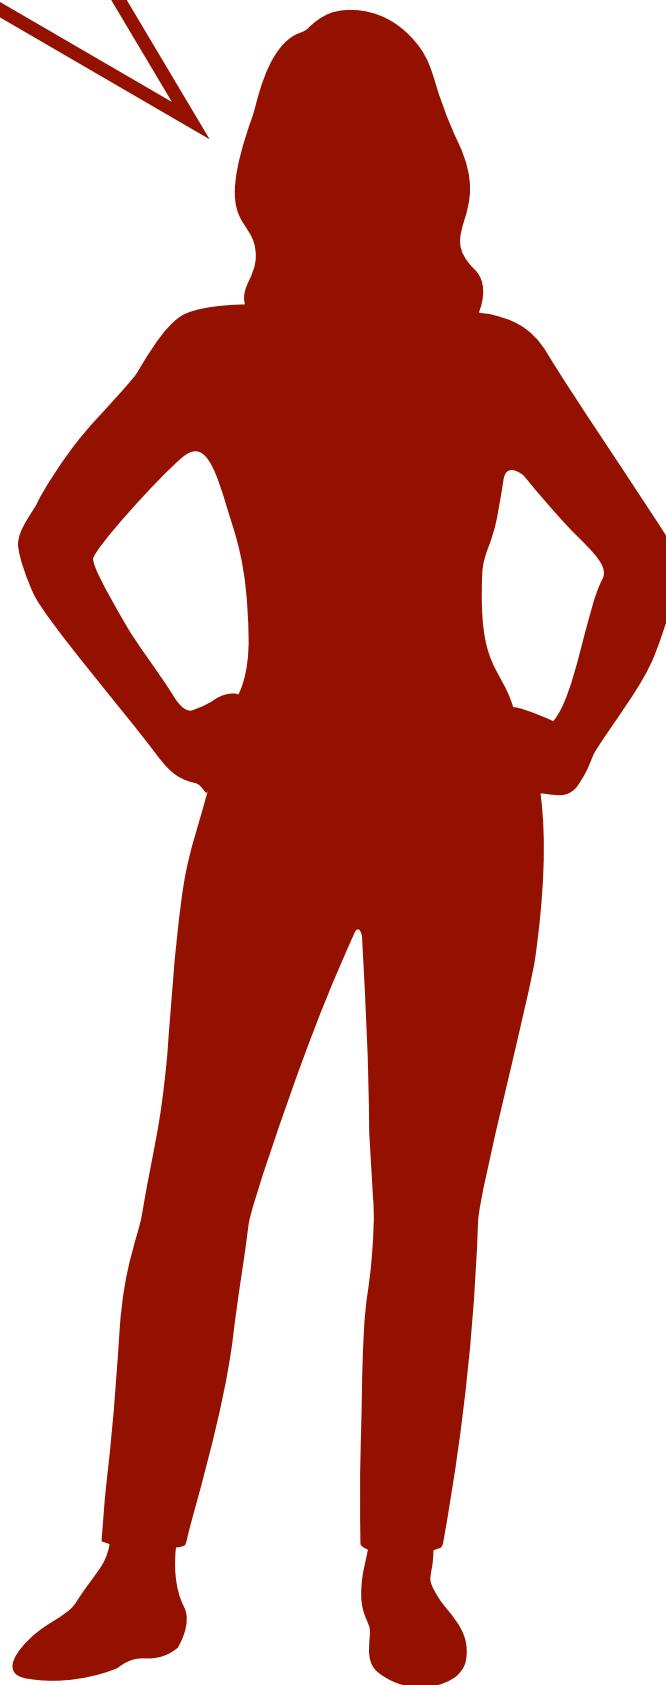


GitHub first



- ▶ Step 1: Create a new repo on GitHub
- ▶ Step 2: Copy the repo URL
- ▶ Step 3: Clone it using RStudio
- ▶ Step 4: Make changes locally
- ▶ Step 6: Commit and push to GitHub
- ▶ Step 7: Confirm your changes have propagated to GitHub

Local first



I have been working on
a project for a while, and now
I'm realising I should have
been tracking it with git.

- ▶ Step 1: Create an RStudio Project from existing directory (if an `.Rproj` file doesn't already exist)
- ▶ Step 2: `usethis::use_git()` and follow instructions
- ▶ Step 3: `usethis::use_github()` and follow instructions

demo

git & github

more resources . . .

- ▶ Learn more about using **Git and GitHub with R**:
 - ▶ Book: happygitwithr.com
- ▶ Learn more about **Git setup**:
 - ▶ Documentation: usethis.r-lib.org/articles/articles/usethis-setup.html

?

automate
your
process



raw-data



processed-data



scripts

- └ 00-analyse.R
- └ 01-load-packages.R
- └ 02-load-data.R
- └ 03-clean-data.R
- └ 04-explore.R
- └ 05-model.R
- └ 06-summarise.R



```
00-analyse.R x
Source on Save | Run | Source | R Script

1 # run all -----
2
3 source("01-load-packages.R")
4 source("02-load-data.R")
5 source("03-clean-data.R")
6 source("04-explore.R")
7 source("05-model.R")
8 source("06-summarise.R")

1:1 # run all
```



figures



manuscript



recommended reading

minimal make A minimal tutorial on make

I would argue that the most important tool for reproducible research is not [Sweave](#) or [knitr](#) but [GNU make](#).

Consider, for example, all of the files associated with a manuscript. In the simplest case, I would have an [R](#) script for each figure plus a [LaTeX](#) file for the main text. And then a [BibTeX](#) file for the references.

Compiling the final PDF is a bit of work:

- Run each R script through R to produce the relevant figure.
- Run latex and then bibtex and then latex a couple of more times.

And the R scripts need to be run before latex is, and only if they've changed.

A simple example

[GNU make](#) makes this easy. In your directory for the manuscript, you create a text file called `Makefile` that looks something like [the following](#) (here using `pdflatex`).

```
mypaper.pdf: mypaper.bib mypaper.tex Figs/fig1.pdf Figs/fig2.pdf  
    pdflatex mypaper  
    bibtex mypaper  
    pdflatex mypaper  
    pdflatex mypaper  
  
Figs/fig1.pdf: R/fig1.R  
    cd R;R CMD BATCH fig1.R  
  
Figs/fig2.pdf: R/fig2.R  
    cd R;R CMD BATCH fig2.R
```

Each batch of lines indicates a file to be created (the *target*), the files it depends on (the *prerequisites*), and then a set of commands needed to construct the target from the dependent files. Note that the lines with the commands *must start with a tab character (not spaces)*.



The targets R Package User Manual

Will Landau

Copyright Eli Lilly and Company

recommended reading

Chapter 1 Introduction

The `targets` package is a [Make](#)-like pipeline toolkit for Statistics and data science in R. With `targets`, you can maintain a reproducible workflow without repeating yourself. `targets` learns how your pipeline fits together, skips costly runtime for tasks that are already up to date, runs only the necessary computation, supports implicit parallel computing, abstracts files as R objects, and shows tangible evidence that the results match the underlying code and data.

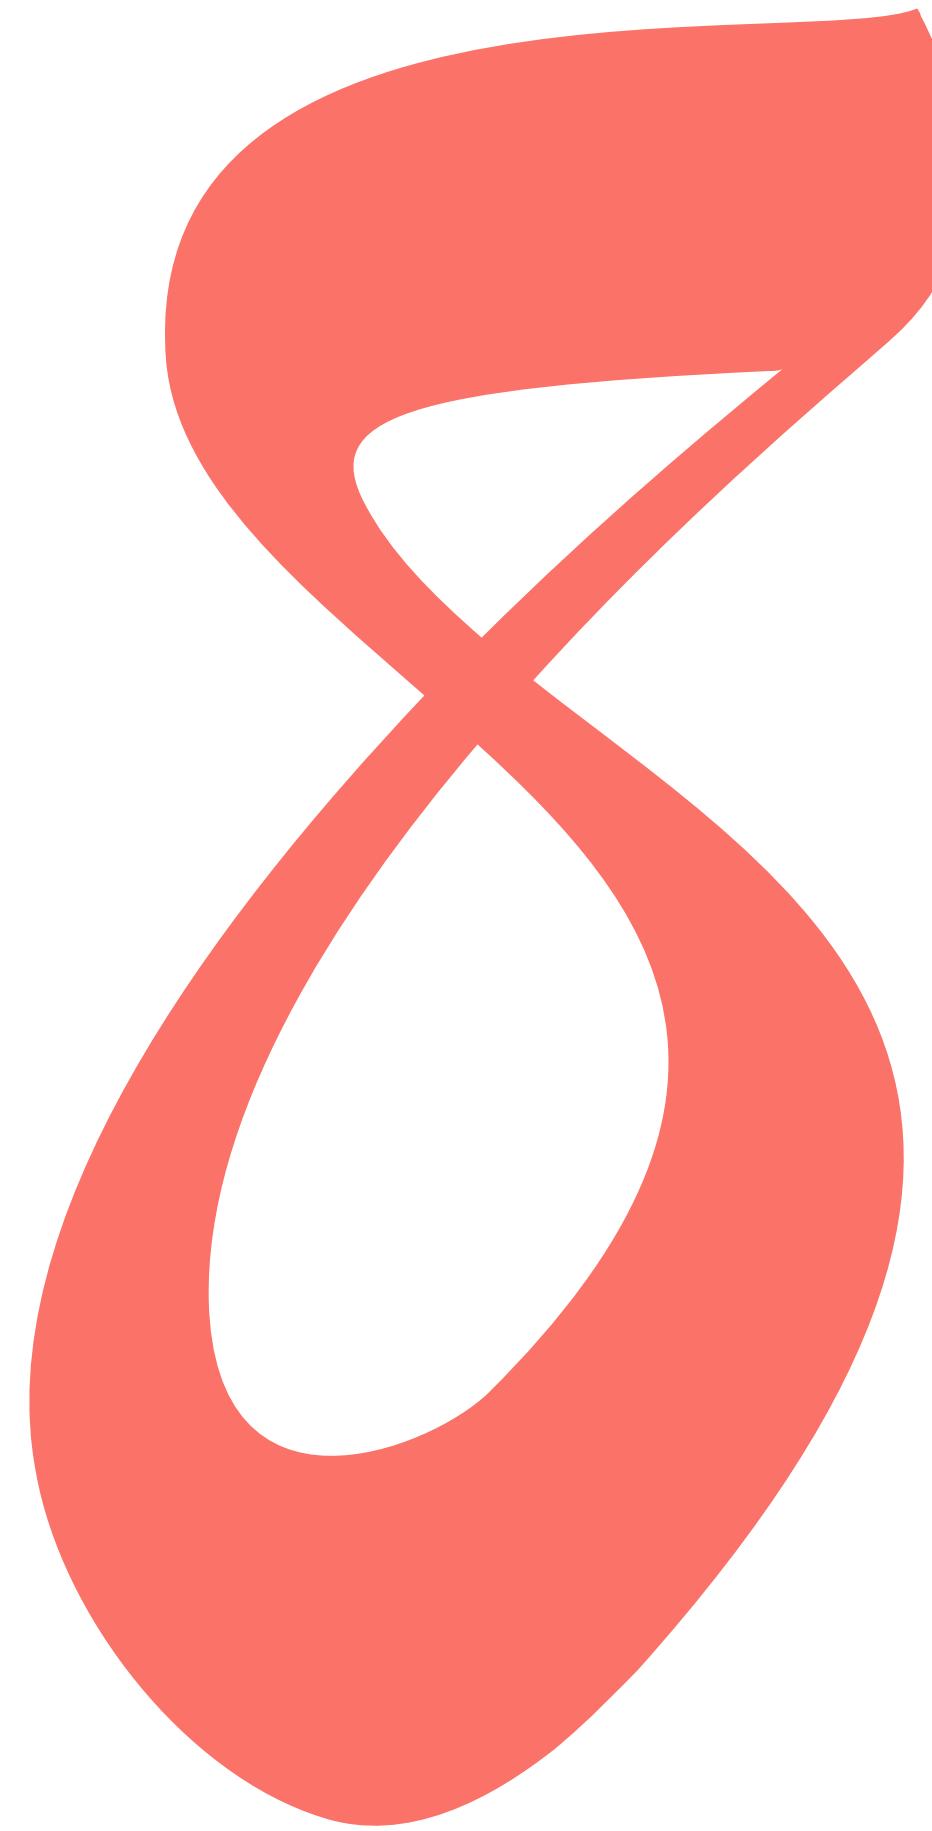
1.1 Motivation

Data analysis can be slow. A round of scientific computation can take several minutes, hours, or even days to complete. After it finishes, if you update your code or data, your hard-earned results may no longer be valid. Unchecked, this invalidation creates chronic [Sisyphean](#) loop:

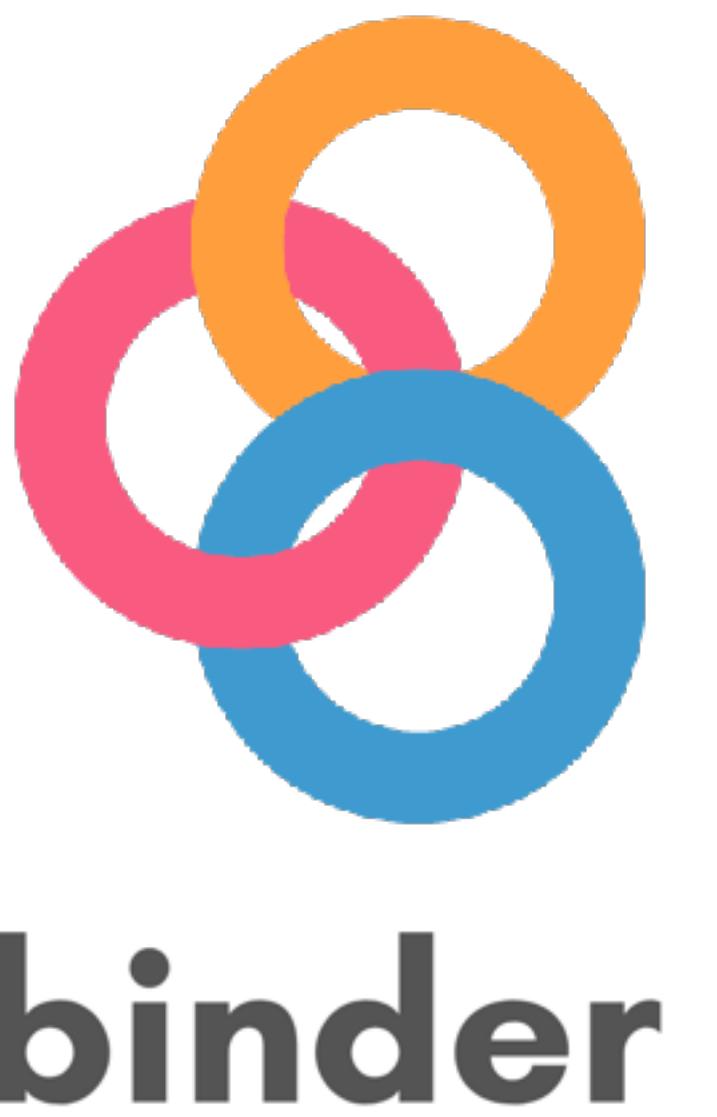
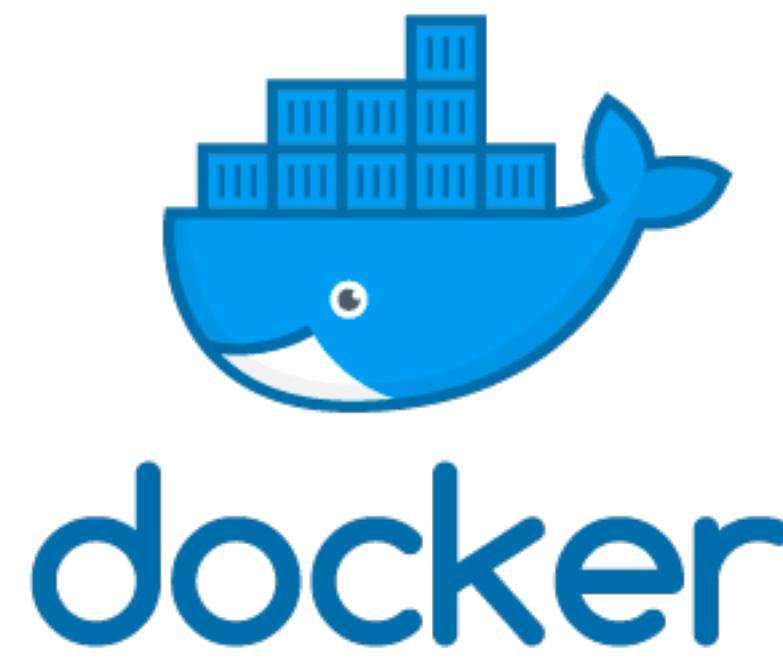
1. Launch the code.
2. Wait while it runs.
3. Discover an issue.
4. Restart from scratch.

1.2 Pipeline toolkits

Pipeline toolkits like [GNU Make](#) break the cycle. They watch the dependency graph of the whole workflow and skip steps, or “targets”, whose code, data, and upstream dependencies have not changed since the last run of the pipeline. When all targets are up to date, this is evidence that the results match the underlying code and data, which helps us trust the results and confirm the computation is reproducible.



share
computing
environment



- 1 organize your project
- 2 write **READMEs** liberally
- 3 keep data **tidy & machine readable**
- 4 comment your code
- 5 use **literate programming**
- 6 use **version control**
- 7 automate your process
- 8 share computing **environment**

PERSPECTIVE

Good enough practices in scientific computing

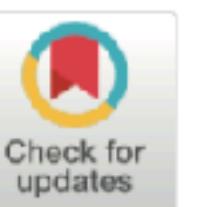
Greg Wilson^{1*}, Jennifer Bryan², Karen Cranston³, Justin Kitzes⁴, Lex Nederbragt⁵, Tracy K. Teal⁶

1 Software Carpentry Foundation, Austin, Texas, United States of America, **2** RStudio and Department of Statistics, University of British Columbia, Vancouver, British Columbia, Canada, **3** Department of Biology, Duke University, Durham, North Carolina, United States of America, **4** Energy and Resources Group, University of California, Berkeley, Berkeley, California, United States of America, **5** Centre for Ecological and Evolutionary Synthesis, University of Oslo, Oslo, Norway, **6** Data Carpentry, Davis, California, United States of America

* These authors contributed equally to this work.

* gwilson@software-carpentry.org

Greg Wilson, Jennifer Bryan, Karen Cranston,
Justin Kitzes, Lex Nederbragt, Tracy K. Teal
“Good enough practices in scientific computing.”
PLoS computational biology 13.6 (2017): e1005510.



OPEN ACCESS

Citation: Wilson G, Bryan J, Cranston K, Kitzes J, Nederbragt L, Teal TK (2017) Good enough practices in scientific computing. PLoS Comput Biol 13(6): e1005510. <https://doi.org/10.1371/journal.pcbi.1005510>

Editor: Francis Ouellette, Ontario Institute for Cancer Research, CANADA

Published: June 22, 2017

Copyright: © 2017 Wilson et al. This is an open access article distributed under the terms of the Creative Commons Attribution License, which permits unrestricted use, distribution, and reproduction in any medium, provided the original author and source are credited.

Funding: The authors received no specific funding for this work.

Competing interests: The authors have declared that no competing interests exist.

Author summary

Computers are now essential in all branches of science, but most researchers are never taught the equivalent of basic lab skills for research computing. As a result, data can get lost, analyses can take much longer than necessary, and researchers are limited in how effectively they can work with software and data. Computing workflows need to follow the same practices as lab projects and notebooks, with organized data, documented steps, and the project structured for reproducibility, but researchers new to computing often don't know where to start. This paper presents a set of good computing practices that every researcher can adopt, regardless of their current level of computational skill. These practices, which encompass data management, programming, collaborating with colleagues, organizing projects, tracking work, and writing manuscripts, are drawn from a wide variety of published sources from our daily lives and from our work with volunteer organizations that have delivered workshops to over 11,000 people since 2010.

Overview

We present a set of computing tools and techniques that every researcher can and should consider adopting. These recommendations synthesize inspiration from our own work, from the experiences of the thousands of people who have taken part in Software Carpentry and Data Carpentry workshops over the past 6 years, and from a variety of other guides. Our recommendations are aimed specifically at people who are new to research computing.

Introduction

Three years ago, a group of researchers involved in Software Carpentry and Data Carpentry wrote a paper called “Best Practices for Scientific Computing” [1]. That paper provided recommendations for people who were already doing significant amounts of computation in their research. However, as computing has become an essential part of science for all researchers, there is a larger group of people new to scientific computing, and the question then becomes, “where to start?”

Improve your workflow for reproducible science



bit.ly/improve-repro-workflow

@minebocek

mine-cetinkaya-rundel

cetinkaya.mine@gmail.com